



Introduction to Kubernetes using AKS

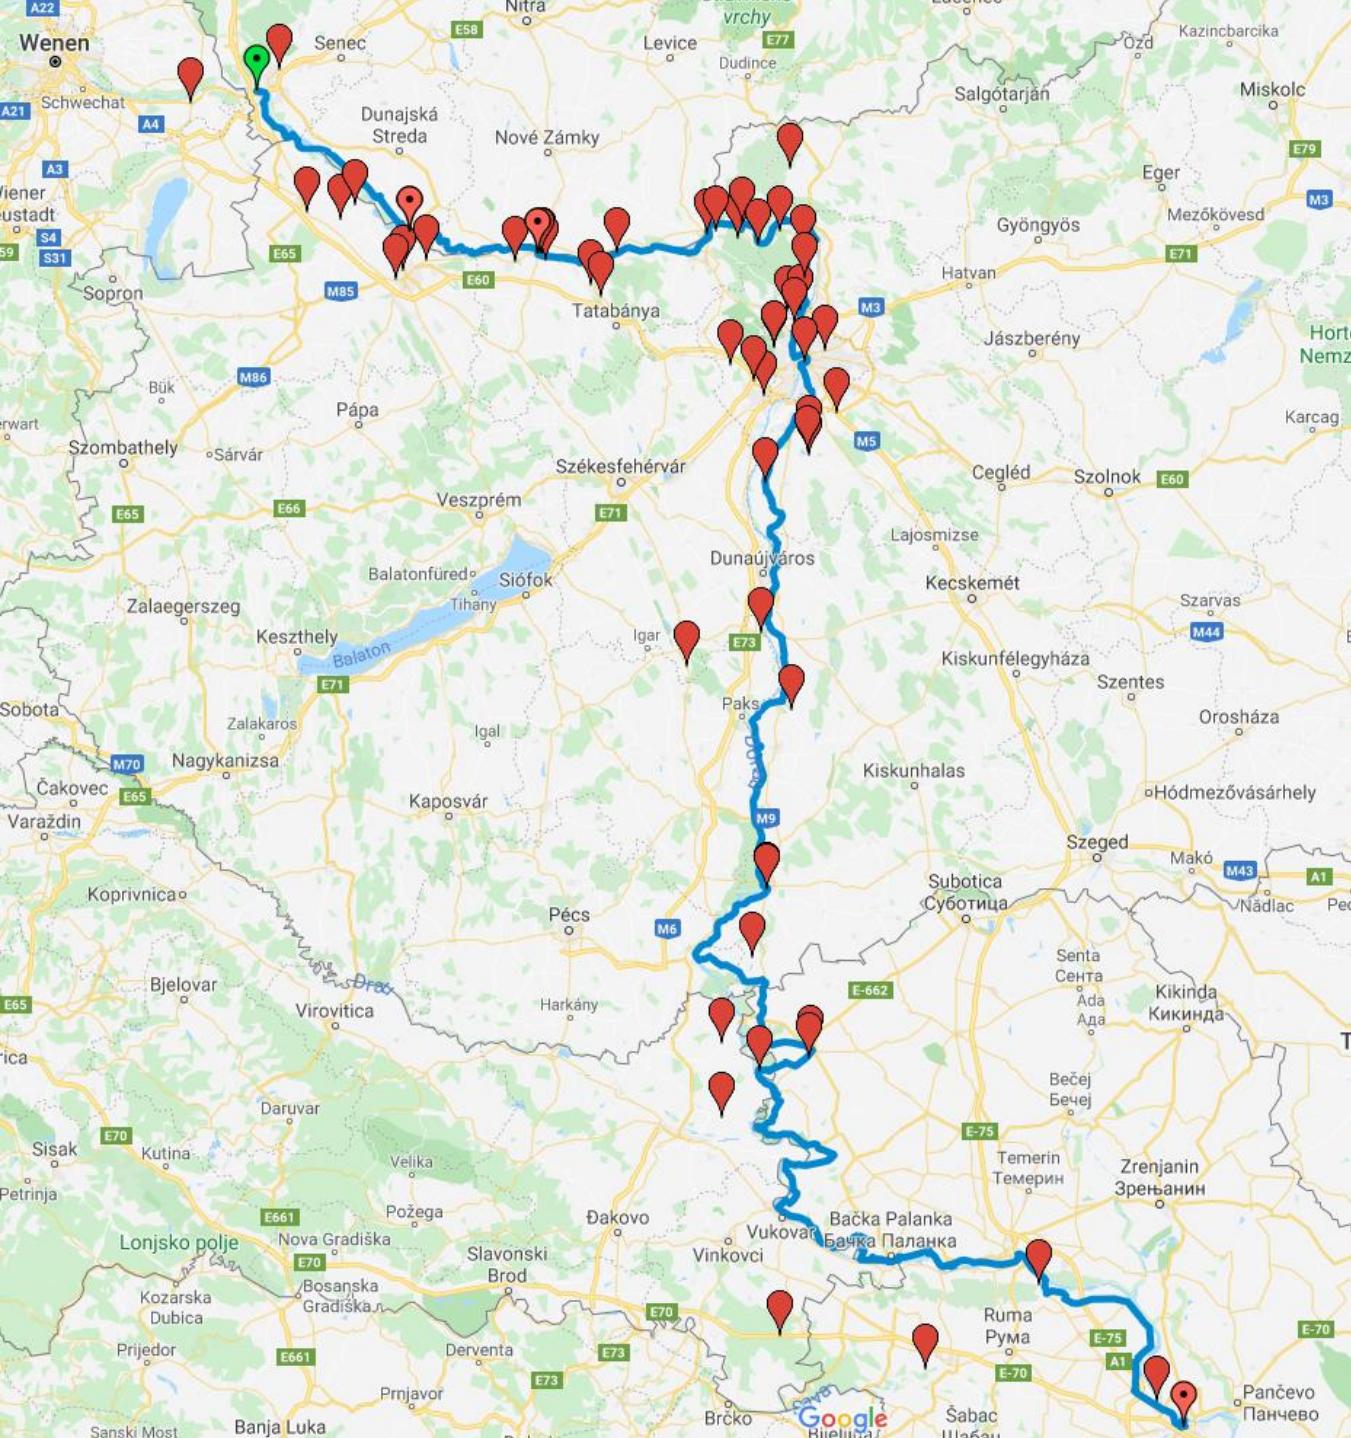
Pascal Naber



MSC HOME TERMINAL



@pascalnaber





kubernetes



docker



@pascalnaber

A photograph of two hands clasped together in the foreground, set against a warm sunset or sunrise sky with orange and blue hues. The hands are dark-skinned.

WHY I LIKE Containers



@pascalnaber

They are FAST





They are **PORTABLE**



@pascalnaber

They are ISOLATED

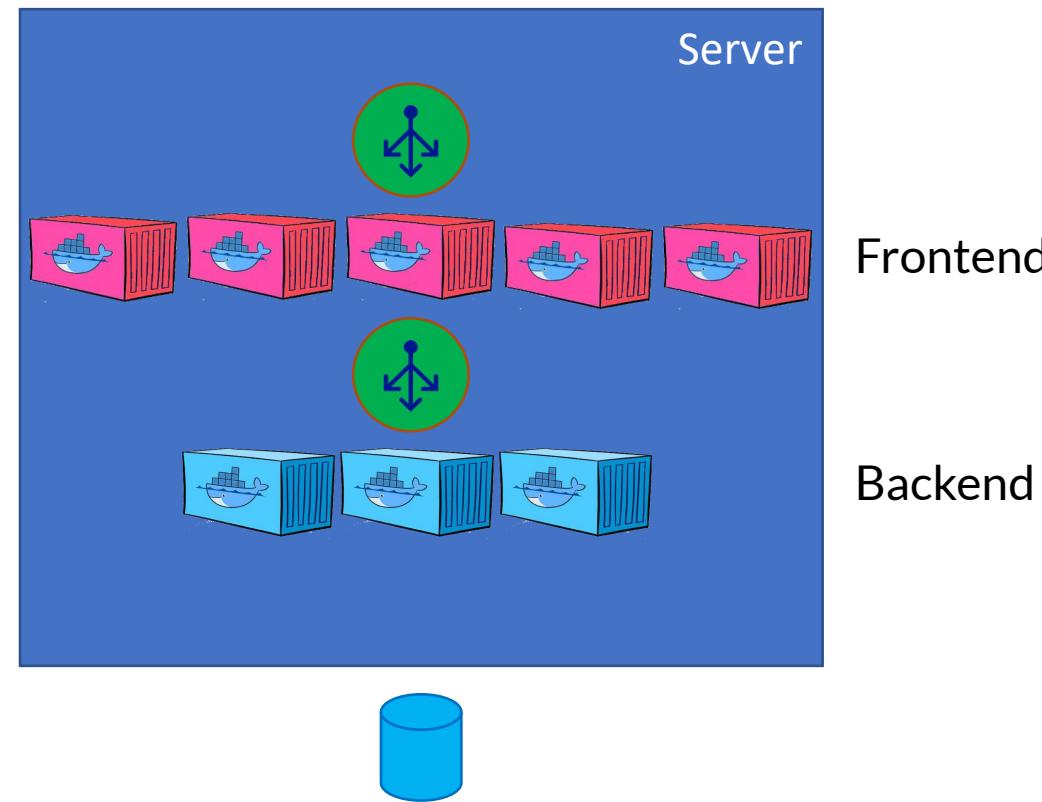


A close-up photograph of a person's hands working on a large sheet of paper. The person is using a white pencil to draw on the paper, which is resting on a wooden surface. A wooden ruler is placed horizontally across the paper, and the person is using it to measure and draw straight lines. In the background, another person is visible, wearing a plaid shirt and holding a red object.

Moving to production

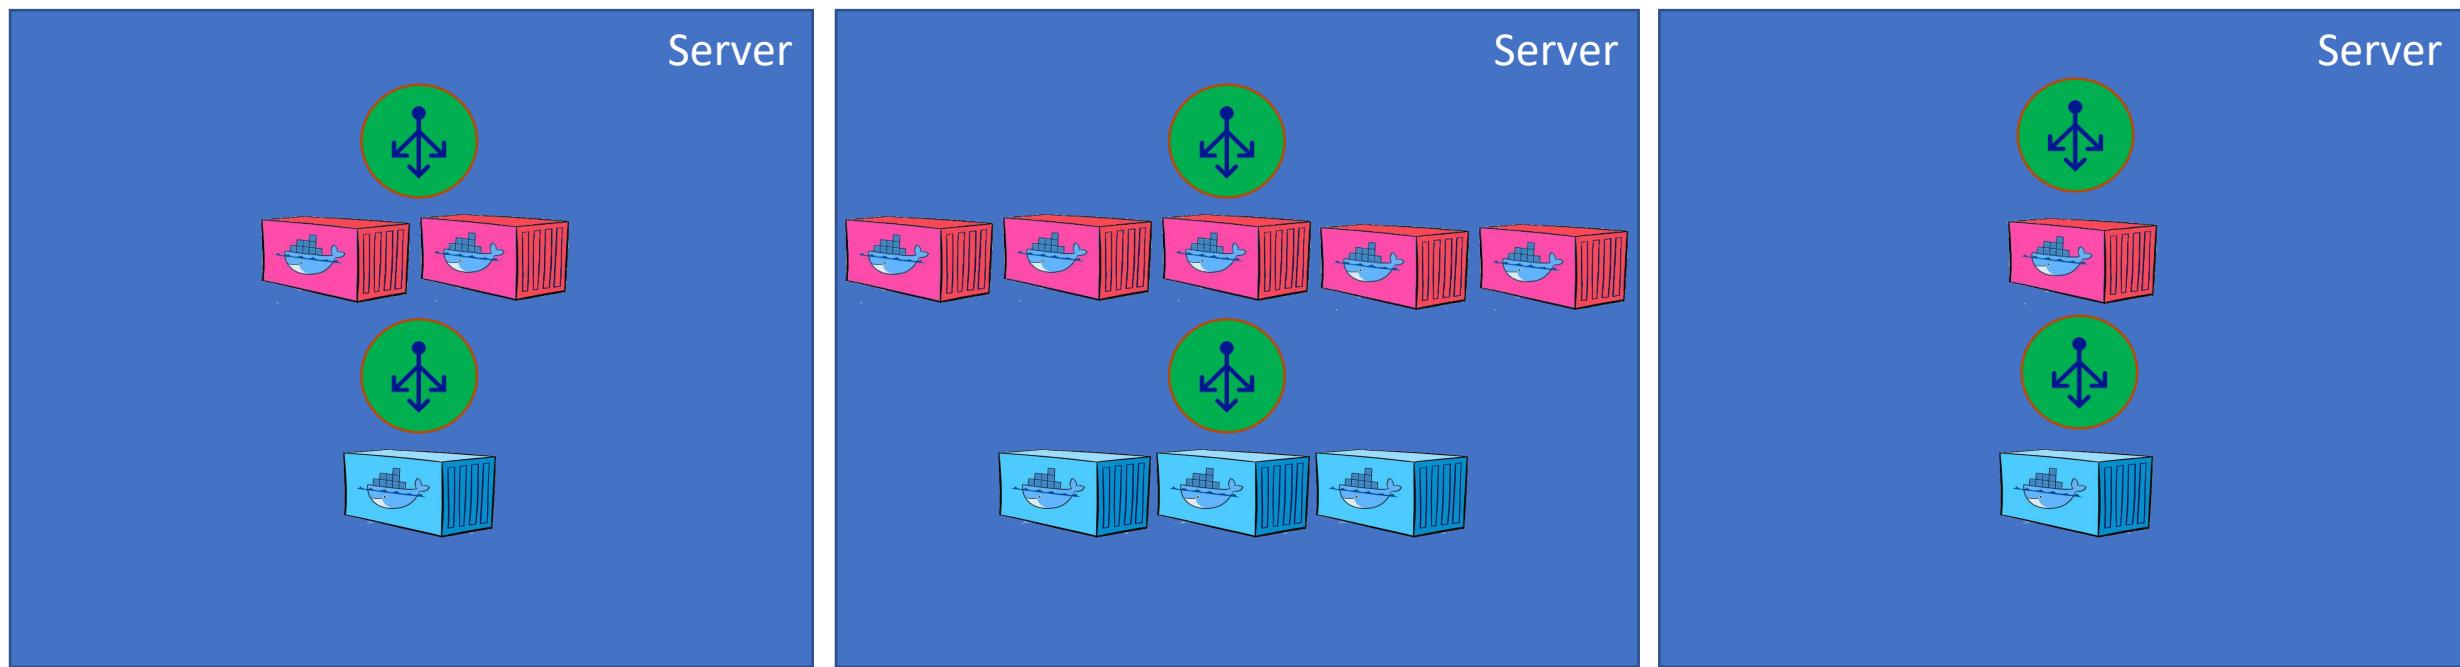
Scaling, load balancing, fault tolerance

Scaling, Load balancing & Fault tolerance



Scaling, Load balancing & Fault tolerance

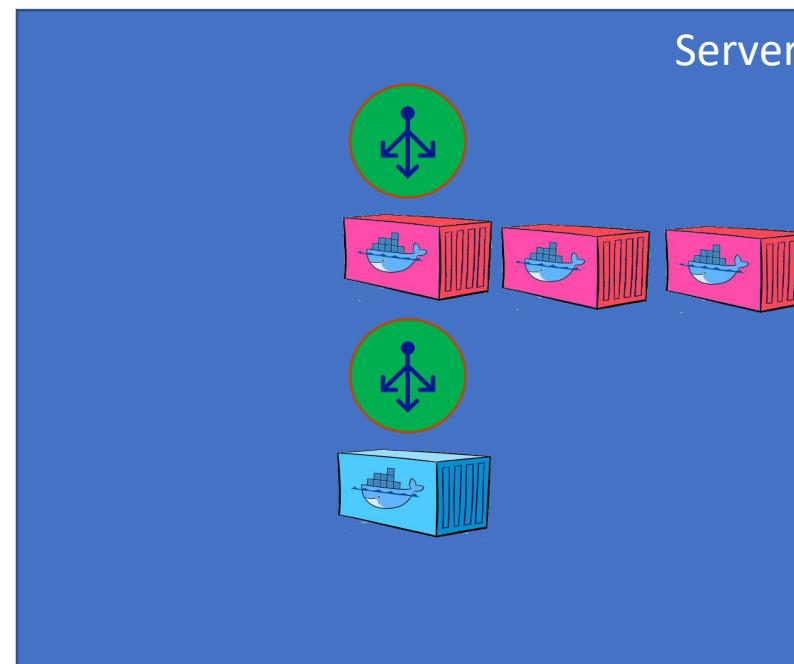
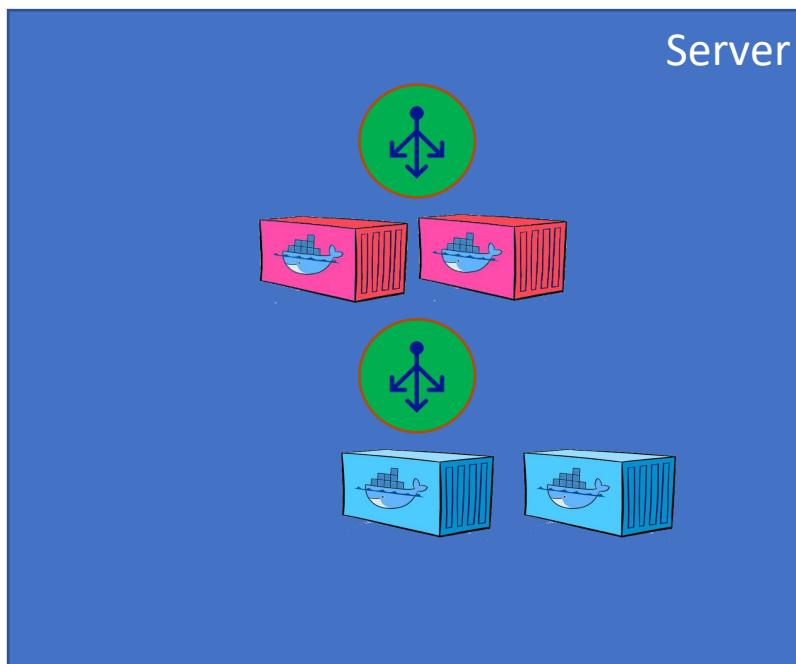
Container orchestration



@pascalnaber

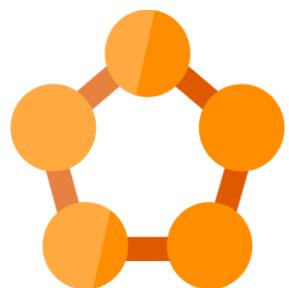
Scaling, Load balancing & Fault tolerance

Container orchestration



@pascalnaber

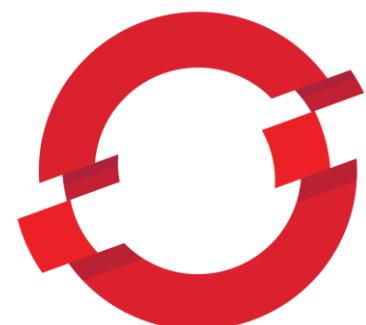
Container orchestrators



Service Fabric



kubernetes



OPENSHIFT

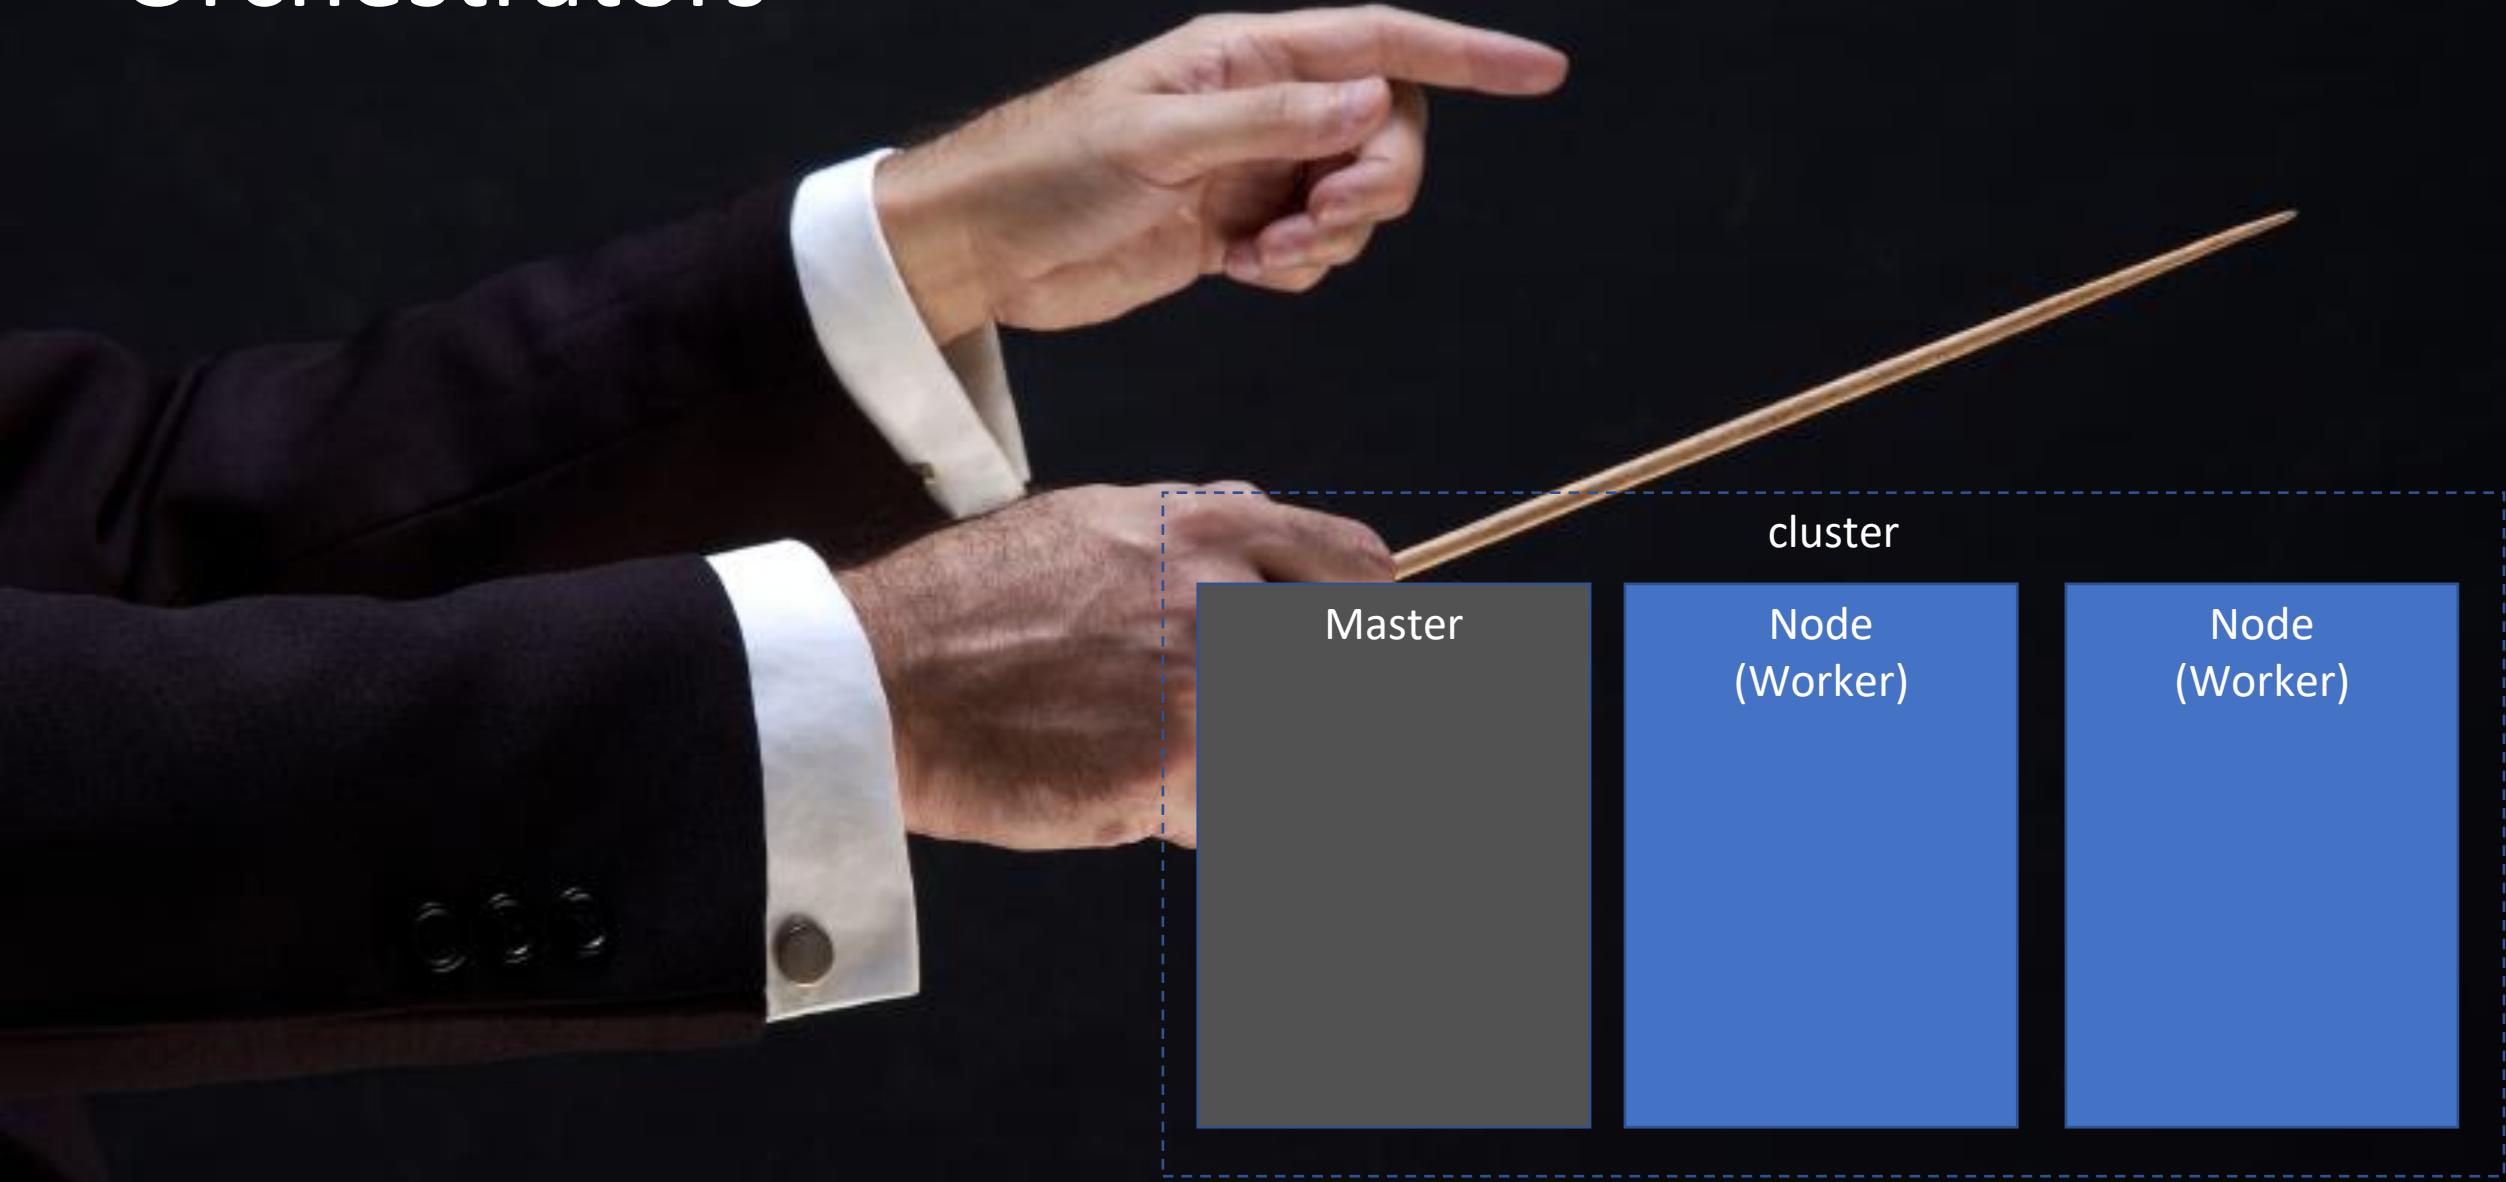


@pascalnaber

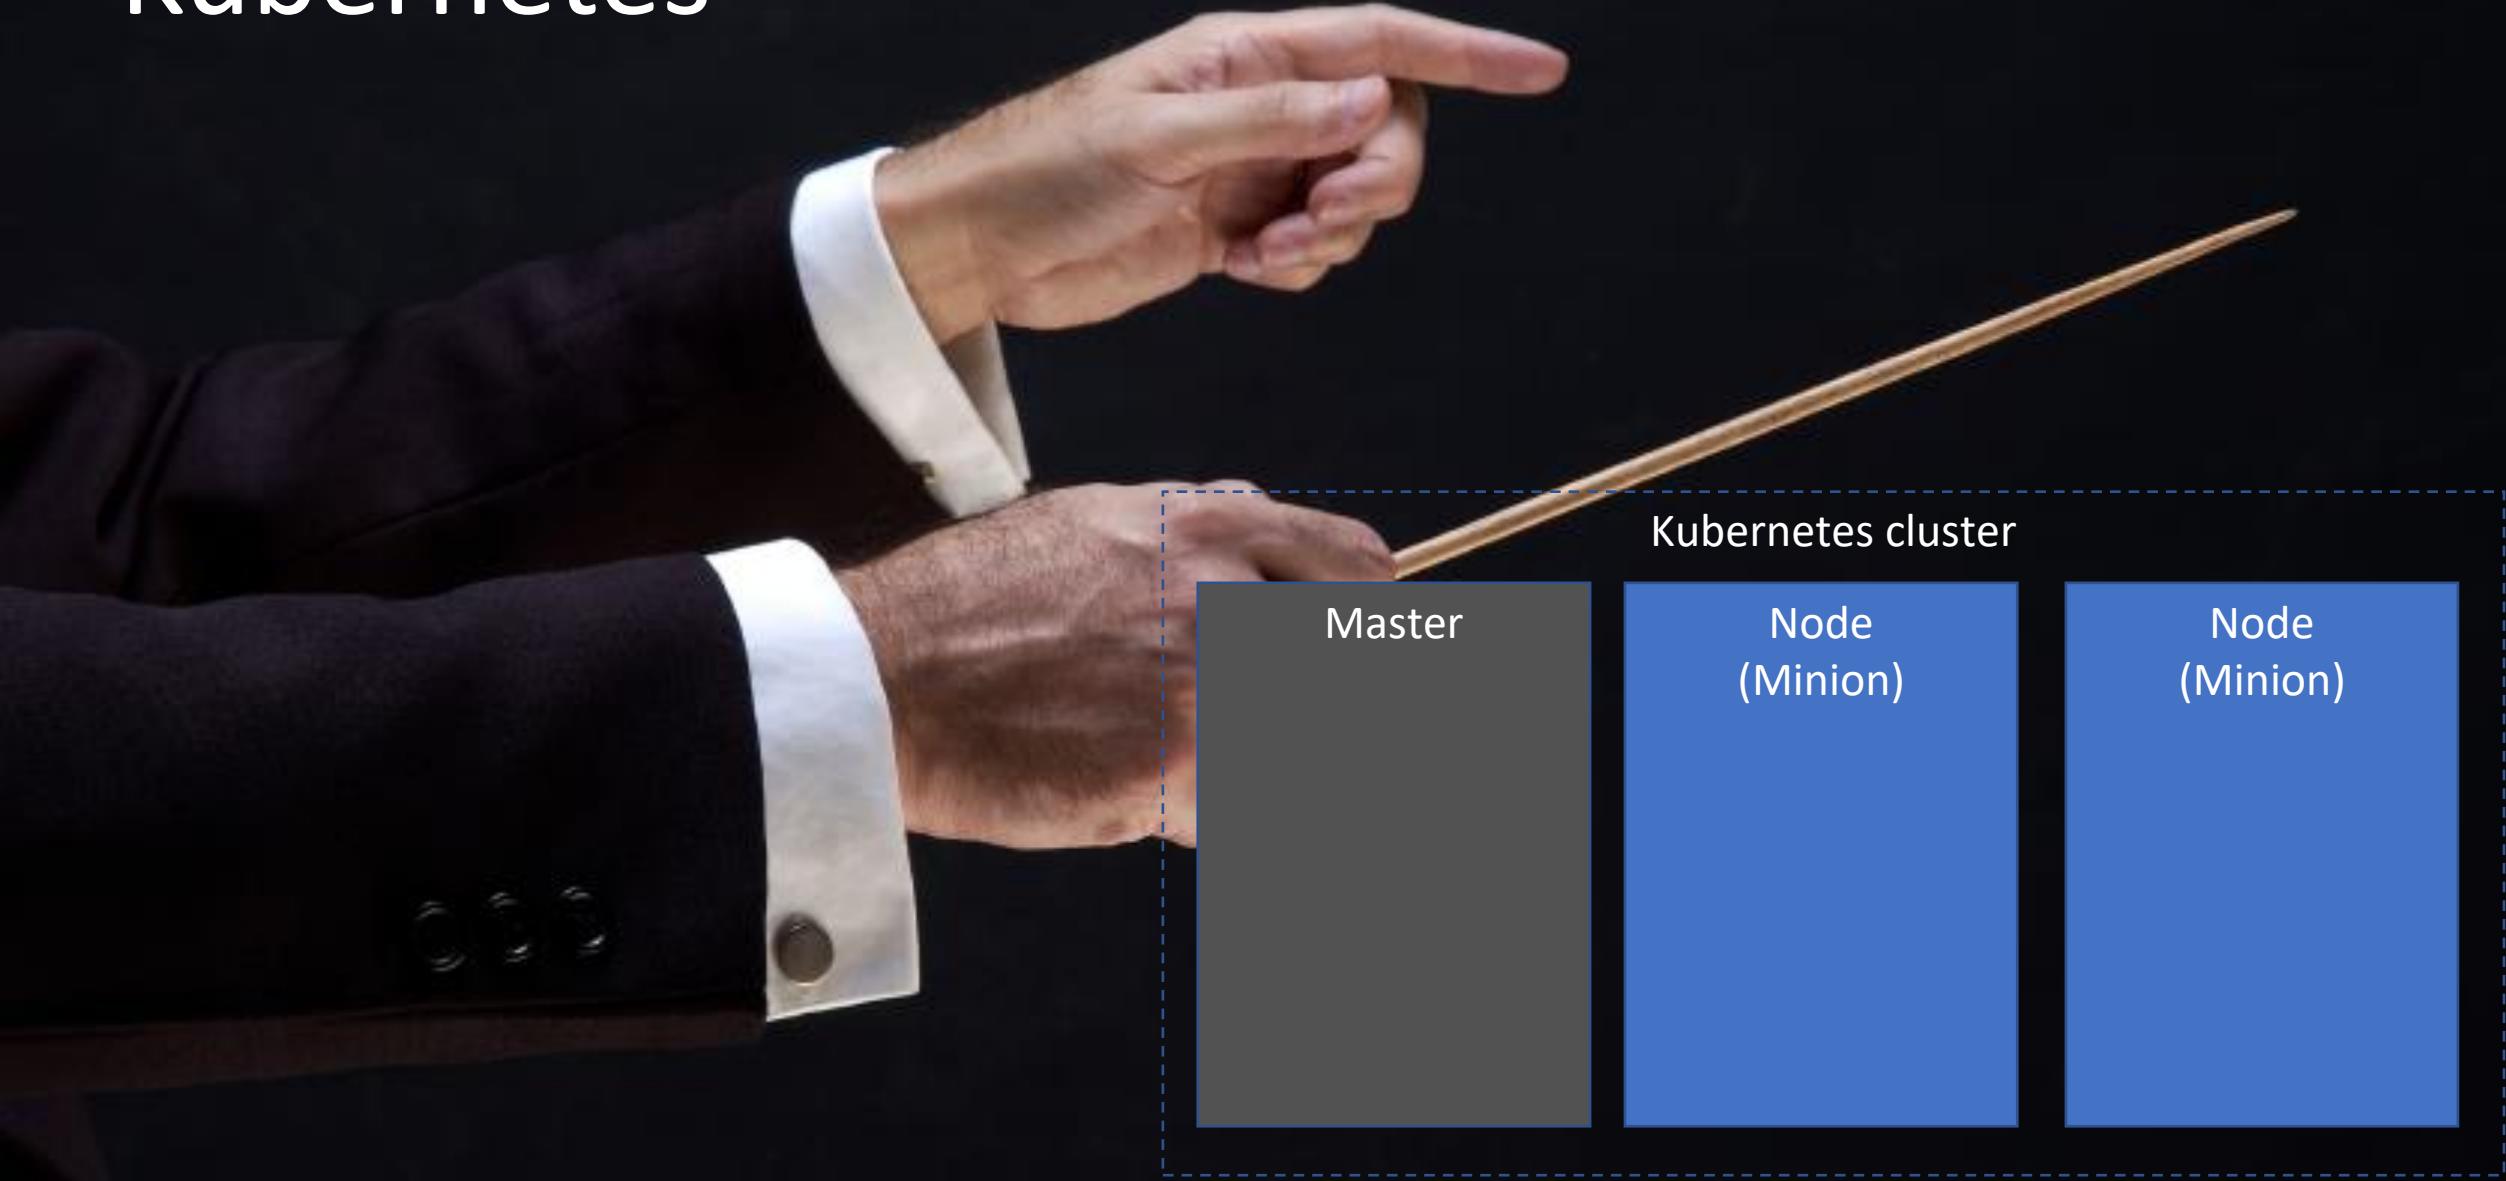
Google Trend



Orchestrators



Kubernetes

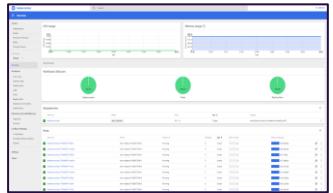


Kubernetes architecture

Kubectl
(CLI)

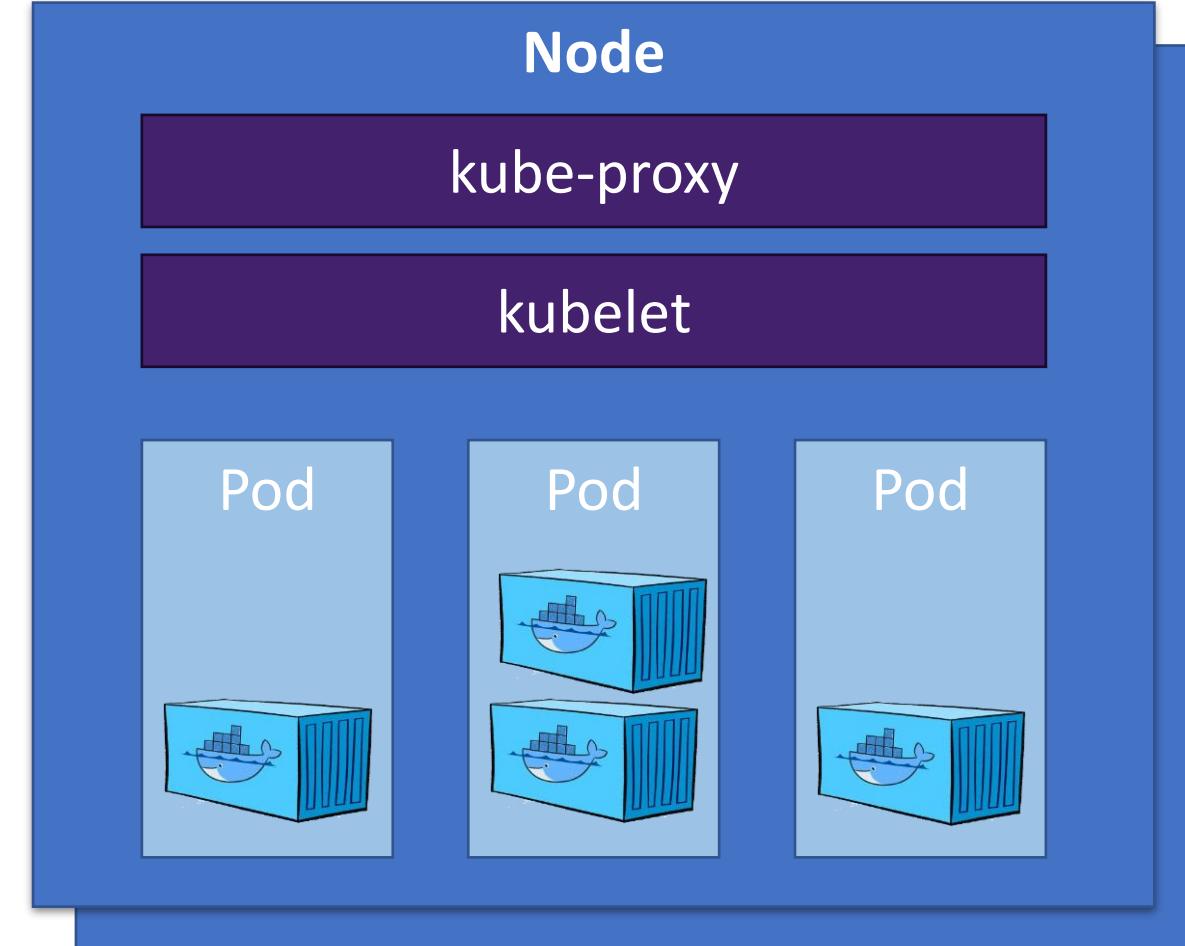
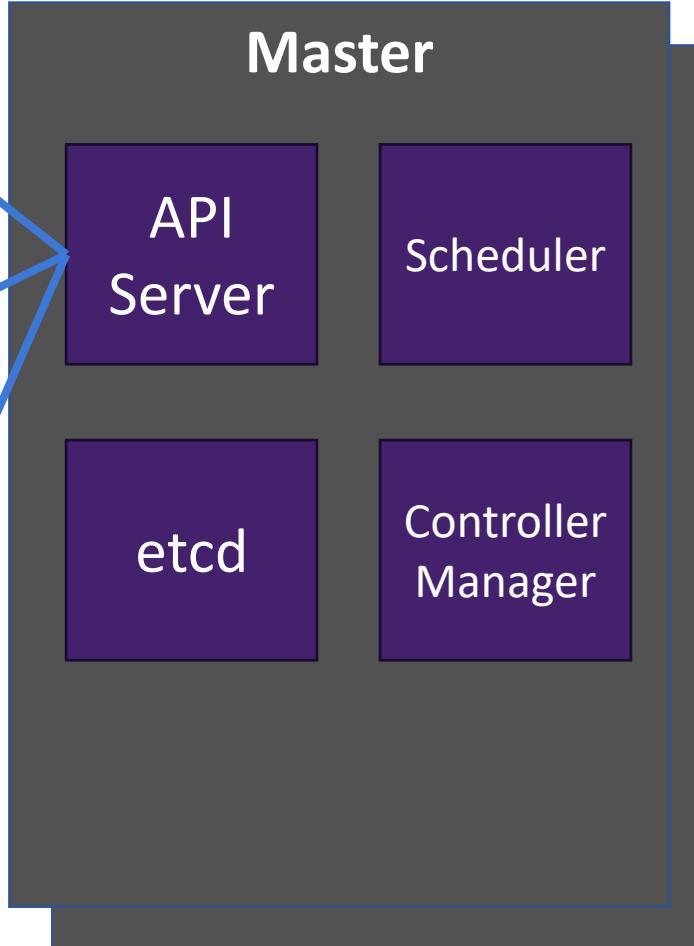
```
C:\>kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
leaderboard-apl-1-75f5d6ff47-4fg8s  1/1     Running   0          2d
leaderboard-apl-1-75f5d6ff47-67h1t  1/1     Running   0          2d
leaderboard-apl-1-75f5d6ff47-jpvu7  1/1     Running   0          2d
leaderboard-apl-1-75f5d6ff47-qwuhd  1/1     Running   0          2d
leaderboard-apl-1-75f5d6ff47-cq9g  1/1     Running   0          2d
leaderboard-apl-1-75f5d6ff47-chnt  1/1     Running   0          2d
leaderboard-apl-1-75f5d6ff47-hxct  1/1     Running   0          2d
leaderboard-apl-1-75f5d6ff47-1z92  1/1     Running   0          2d
leaderboard-apl-1-75f5d6ff47-1zbf  1/1     Running   0          2d
leaderboard-apl-1-75f5d6ff47-ss22m 1/1     Running   0          2d
```

Dashboard
(UI)



REST client
(Code)

{REST}



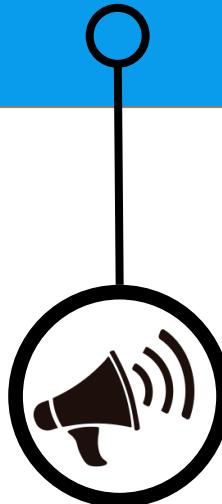
History of Kubernetes



First commit

September 2014

June 2014



Announced

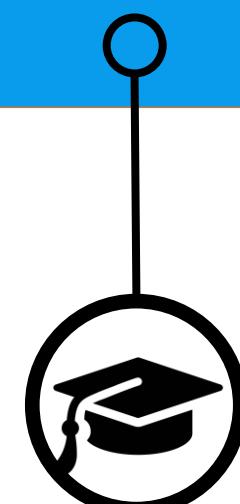


Kubernetes v1

July 2015



March 2018



CNCF 1st
Graduate

Google

CLOUD NATIVE
COMPUTING FOUNDATION

@pascalnaber

Kubernetes

1 2 3 4 5 6 7 8

k8s

OOS

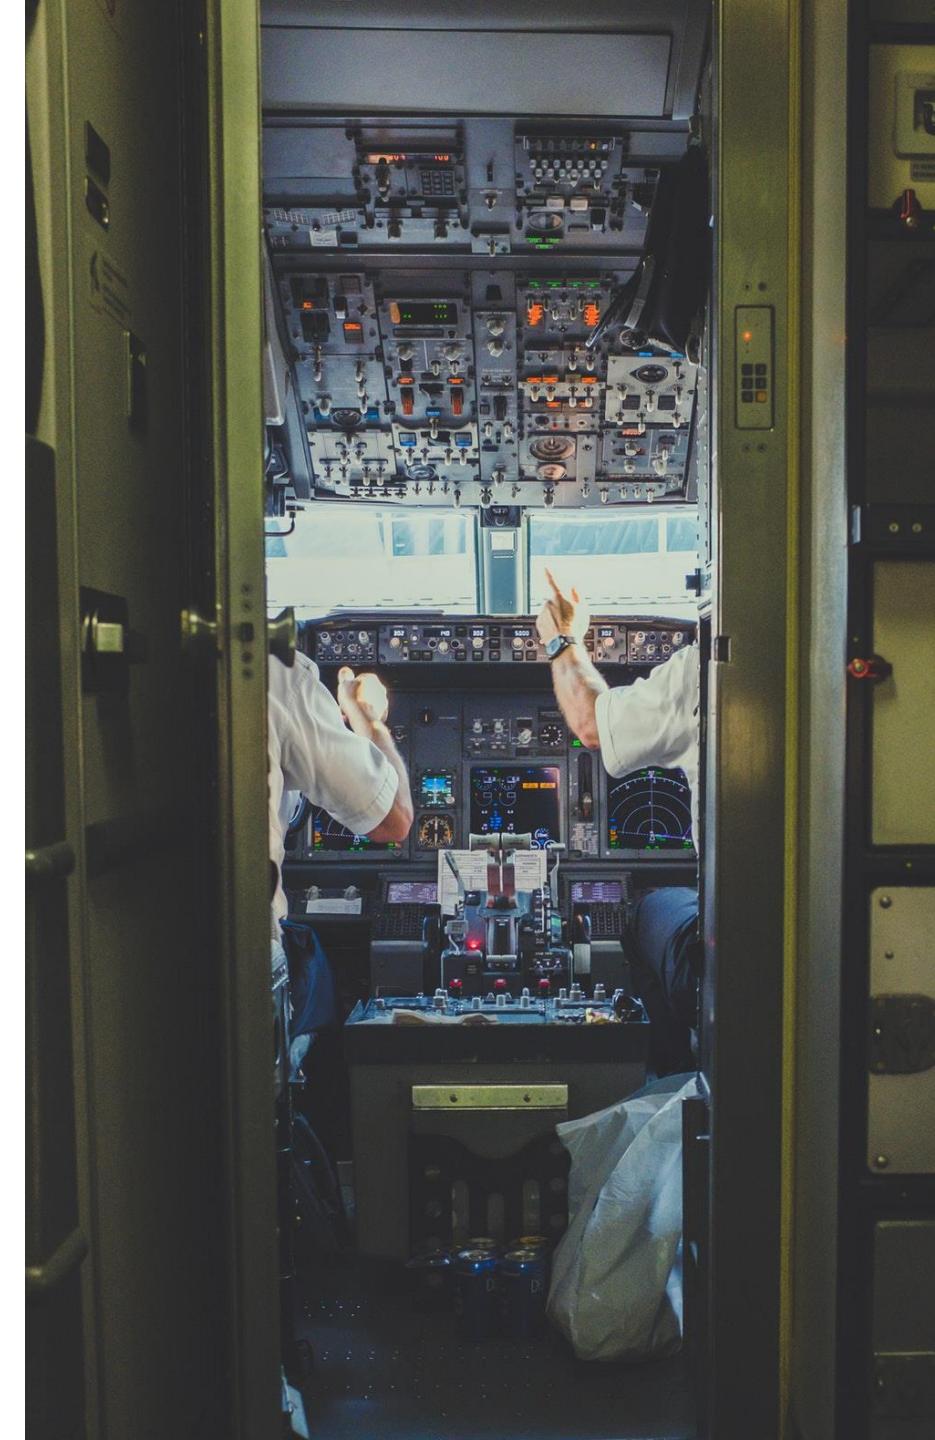


Extensible

Large, rapidly growing ecosystem

Facilitates declarative configuration and automation

*Κυβερνήτης -- Greek:
meaning helmsman or pilot*



Hosting of Kubernetes

Local

Minikube

Docker for Windows

Docker for Mac

Raspberry Pi

Private datacenter

Public Cloud



Google Cloud



GKE (26-8-2015)

EKS (5-6-2018)



ACS (19-4-2016)
AKS (13-6-2018)



@pascalnaber

Azure Kubernetes Service (AKS)



100% managed by Microsoft

€ 0



master



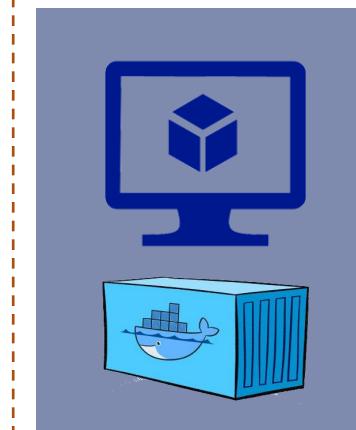
master



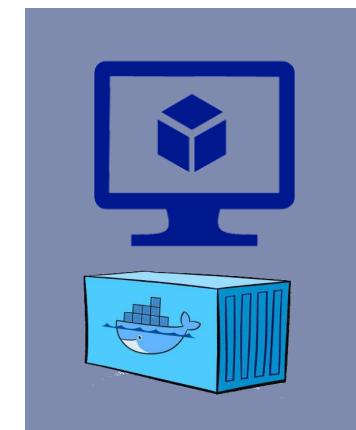
master

IaaS managed by Microsoft

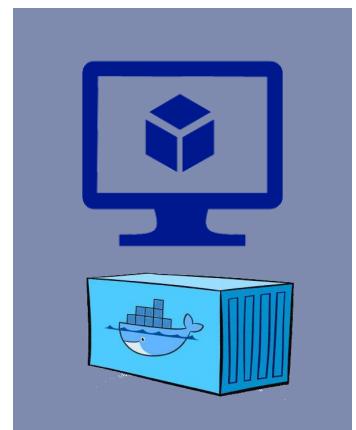
€ ... (VM pricing)



worker



worker

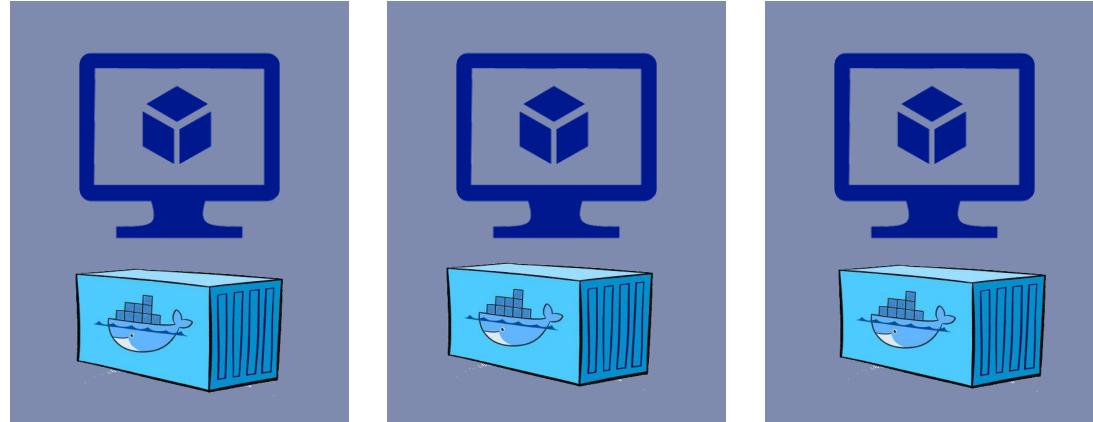


worker

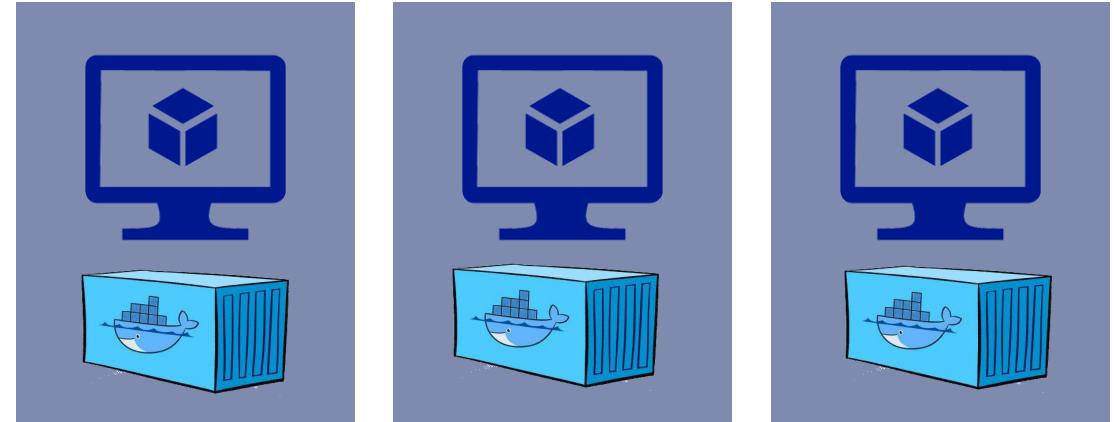


Nodepools

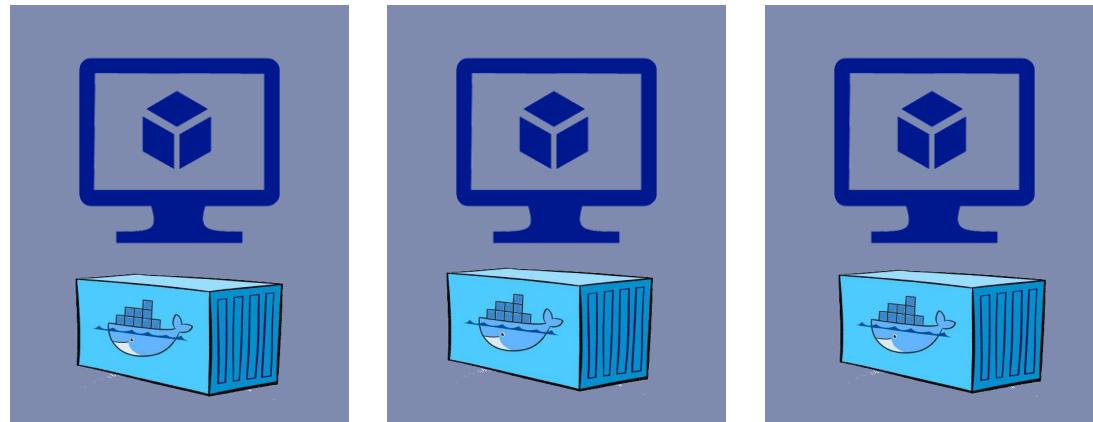
NodePool: Linux | k8s 1.15.7
Standard_DS2_v2



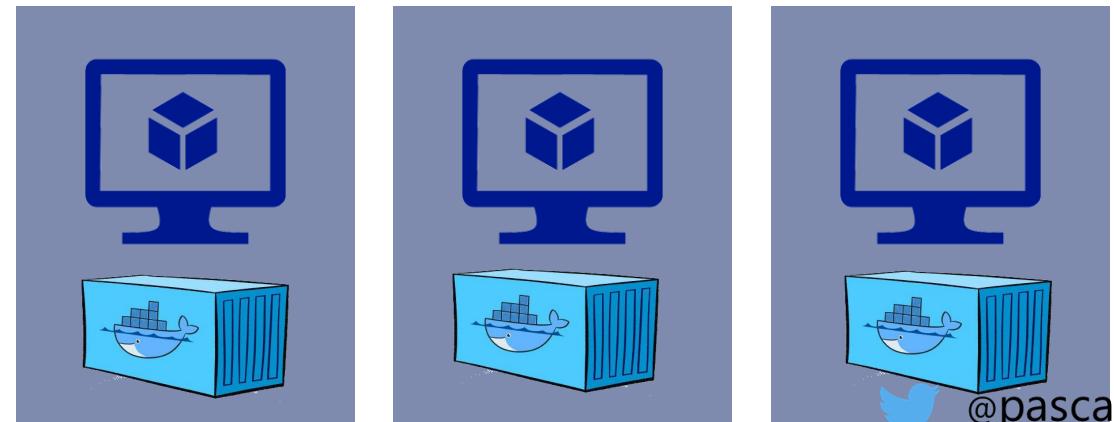
NodePool: Linux | k8s 1.16.7
Standard_NC6 (GPU optimized)



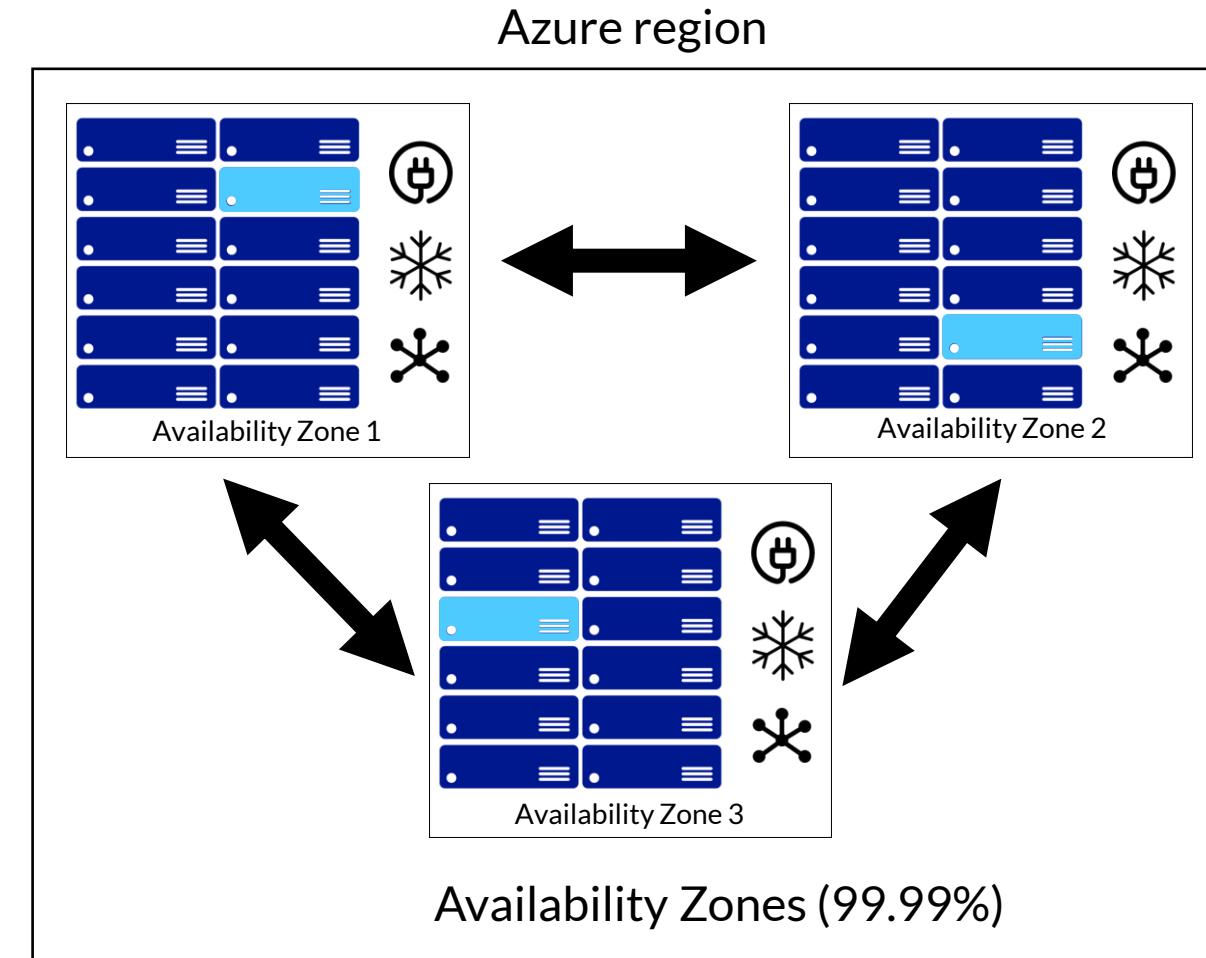
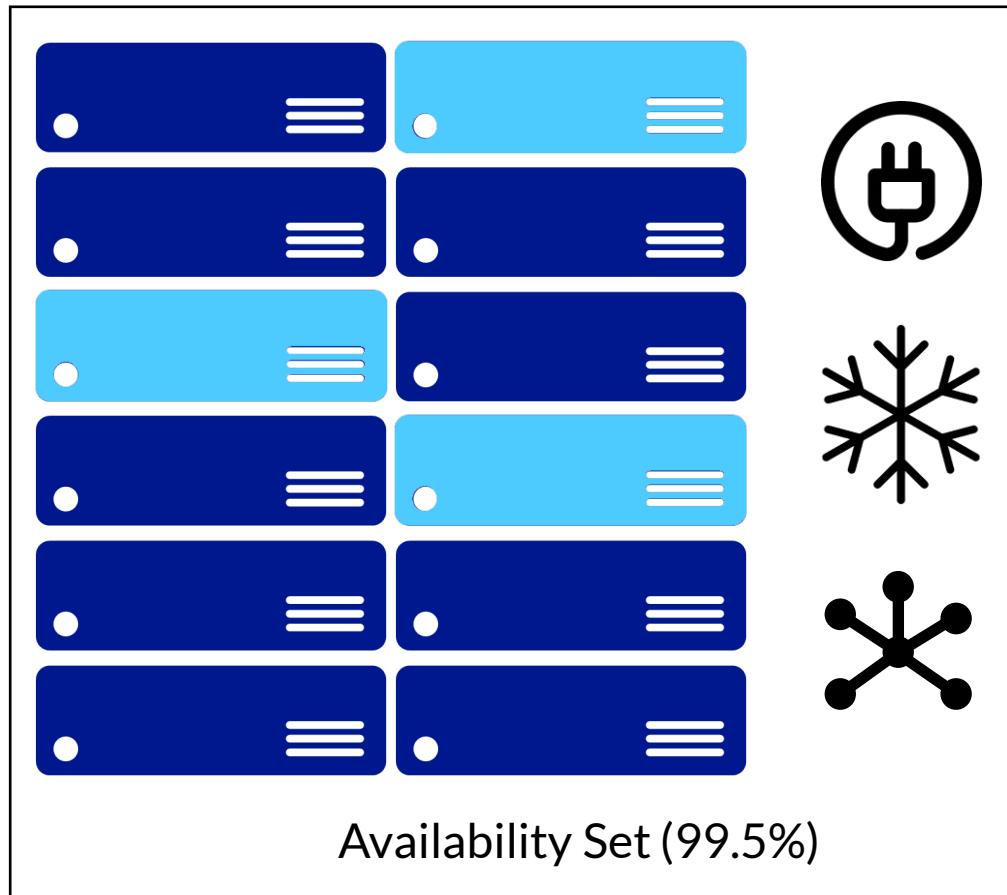
NodePool: Linux | k8s 1.16.7
Standard_DS2_v2



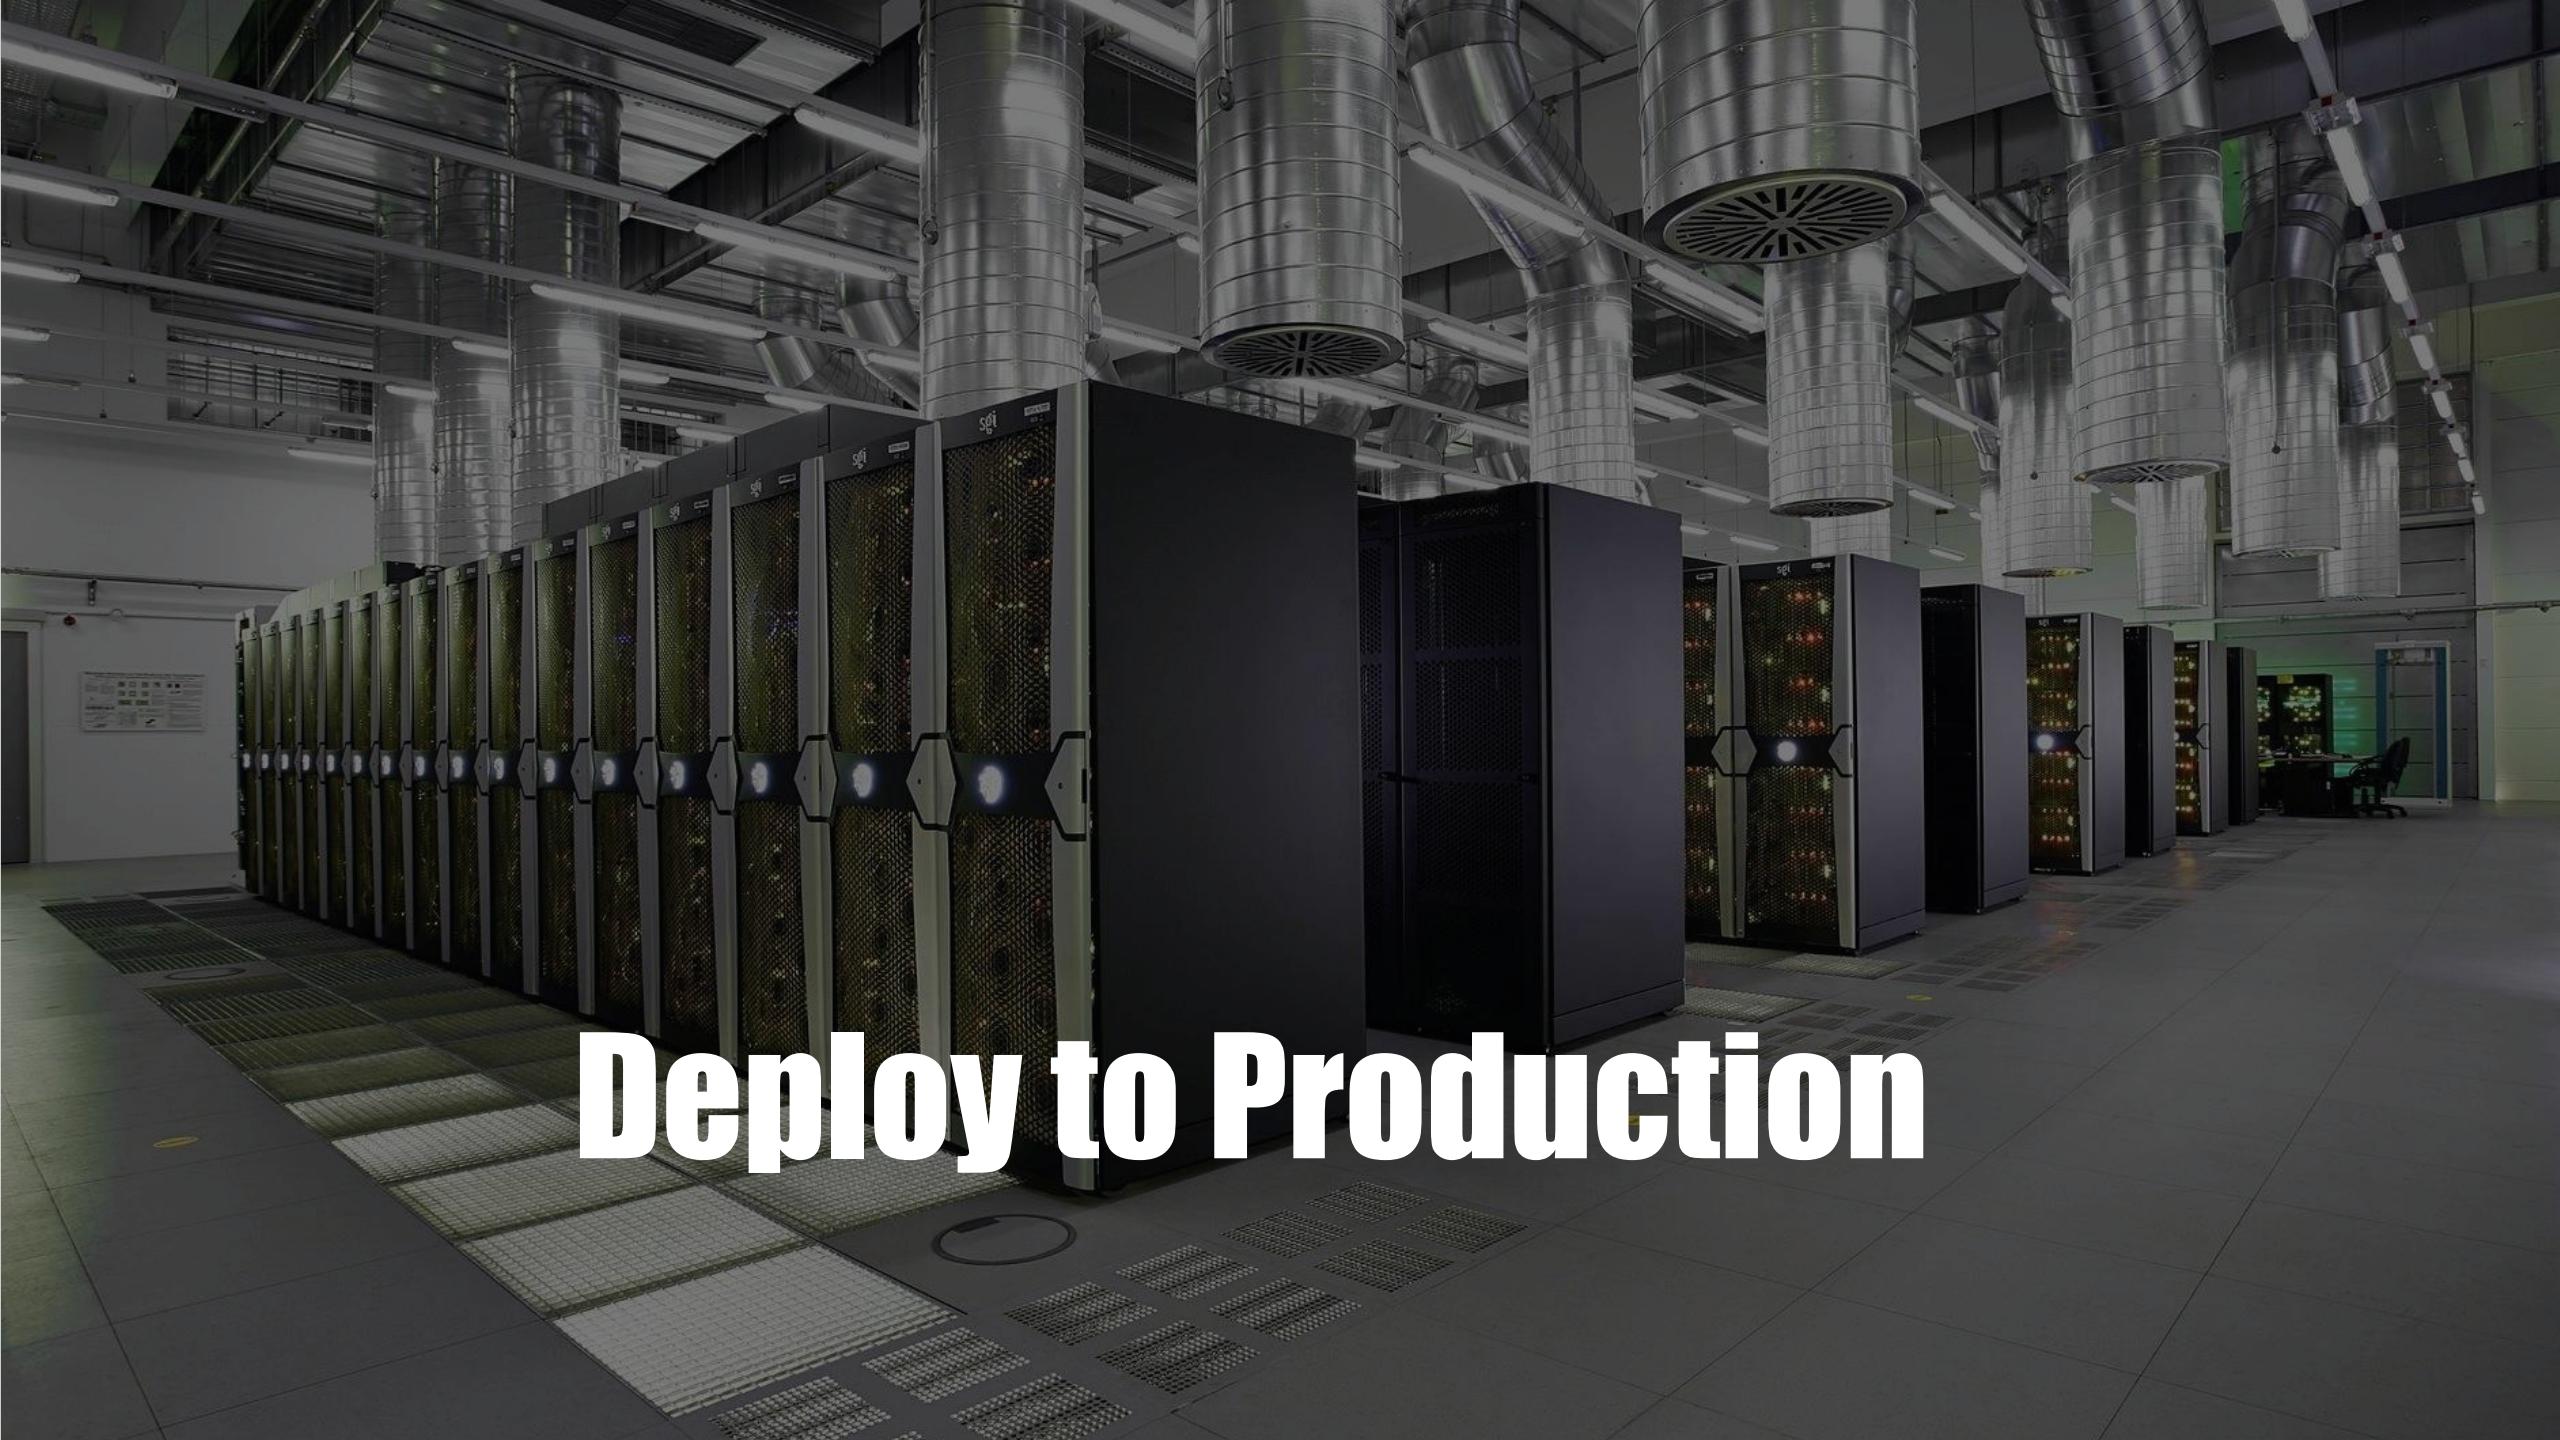
NodePool: Windows | k8s 1.16.7
Standard_DS2_v3



Availability zones



€0.009 per GB inbound & outbound traffic

A photograph of a large server room. In the foreground, there are several tall, dark server racks arranged in a row. The floor has a patterned tile design. Above the racks, a complex network of large, cylindrical metal ducts hangs from the ceiling, used for air conditioning and ventilation. The room is dimly lit, with some light coming from the server racks and the ductwork.

Deploy to Production

Install AKS

ARM Portal

```

"resources": [
    {
        "apiVersion": "2017-08-31",
        "type": "Microsoft.ContainerService/managedClusters",
        "location": "westeurope",
        "name": "[variables('clusterName')]",
        "tags": {
            "department": "[parameters('metadata').department]",
            "projectName": "[parameters('metadata').projectName]",
            "owner": "[parameters('metadata').owner]",
            "environment": "[parameters('metadata').environment]"
        },
        "properties": {
            "kubernetesVersion": "[variables('kubernetesVersion')]",
            "dnsPrefix": "[variables('dnsPrefix')]",
            "agentPoolProfiles": [
                {
                    "name": "agentpool",
                    "osDiskSizeGB": "[variables('osDiskSizeGB')]",
                    "count": "[variables('agentCount')]",
                    "vmSize": "[variables('agentVMSize')]",
                    "osType": "[variables('osType')]",
                    "storageProfile": "ManagedDisks"
                }
            ],
            "linuxProfile": {
                "adminUsername": "[variables('adminUsername')]",
                "ssh": {
                    "publicKeys": [
                        {
                            "keyData": "[variables('sshRSAPublicKey')]"
                        }
                    ]
                }
            },
            "servicePrincipalProfile": {
                "ClientId": "[parameters('servicePrincipalClientId')]",
                "Secret": "[parameters('servicePrincipalClientSecret')]"
            }
        }
    }
]

```

Create Kubernetes cluster

[Basics](#) [Networking](#) [Monitoring](#) [Tags](#) [Review + create](#)

Azure Kubernetes Service (AKS) manages your hosted Kubernetes environment, making it quick and easy to deploy and manage containerized applications without container orchestration expertise. It also eliminates the burden of ongoing operations and maintenance by provisioning, upgrading, and scaling resources on demand, without taking your applications offline. [Learn more about Azure Kubernetes Service](#)

PROJECT DETAILS

Select a subscription to manage deployed resources and costs. Use resource groups like folders to organize and manage all your resources.

* Subscription: Create new Use existing

* Resource group:

CLUSTER DETAILS

* Kubernetes cluster name:

* Region:

* Kubernetes version:

* DNS name prefix:

AUTHENTICATION

* Service principal: [Config my service principal](#)

SCALE

The number and size of nodes in your cluster. For production workloads, at least 3 nodes are recommended for resiliency. For development or test workloads, only one node is required. You will not be able to change the node size after cluster creation, but you will be able to change the number of nodes in your cluster after creation. [Learn more about scaling in Azure Kubernetes Service](#)

* Node size:

* Node count:

```

- name: Create Azure Kubernetes Service
  hosts: localhost
  connection: local
  vars:
    resource_group: myResourceGroup
    location: eastus
    aks_name: myAKScluster
    username: azureuser
    ssh_key: "your_ssh_key"
    client_id: "your_client_id"
    client_secret: "your_client_secret"
  tasks:
    - name: Create resource group
      azure_rm_resourcegroup:
        name: "{{ resource_group }}"
        location: "{{ location }}"
    - name: Create a managed Azure Container Service
      azure_rm_aks:
        name: "{{ aks_name }}"
        location: "{{ location }}"
        resource_group: "{{ resource_group }}"
        dns_prefix: "{{ aks_name }}"
        linux_profile:
          admin_username: "{{ aks_name }}"
          ssh_key: "{{ ssh_key }}"
        service_principal:
          client_id: "{{ client_id }}"
          client_secret: "{{ client_secret }}"
        agent_pool_profiles:
          - name: default
            count: 2
            vm_size: Standard_D2_v2
        tags:
          Environment: Production

```



```

resource "azurerm_resource_group" "k8s" {
  name   = "${var.resource_group_name}"
  location = "${var.location}"
}

resource "azurerm_kubernetes_cluster" "k8s" {
  name           = "${var.cluster_name}"
  location       = "${azurerm_resource_group.k8s.location}"
  resource_group_name = "${azurerm_resource_group.k8s.name}"
  dns_prefix     = "${var.dns_prefix}"

  linux_profile {
    admin_username = "ubuntu"
    ssh_key {
      key_data = "${file("${var.ssh_public_key}")}"
    }
  }

  agent_pool_profile {
    name           = "default"
    count          = "${var.agent_count}"
    vm_size        = "Standard_D2"
    os_type        = "Linux"
    os_disk_size_gb = 30
  }

  service_principal {
    client_id      = "${var.client_id}"
    client_secret  = "${var.client_secret}"
  }

  tags {
    Environment = "Development"
  }
}

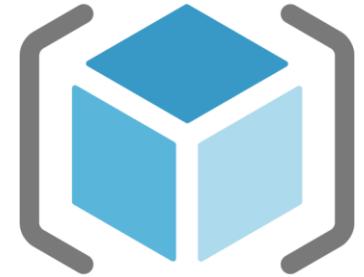
```

```
## Create Kubernetes cluster
az aks create --resource-group communitylive
  --name aksdemo
  --enable-managed-identity
  --kubernetes-version 1.16.7
  --node-vm-size Standard_DS2_v2
  --node-count 3
  --enable-addons monitoring
  --max-pods $AKS_MAX_PODS_PER_NODE
  --ssh-key-value "C:\repos\pascal\ssh\public.pub"
  --load-balancer-sku standard
  --network-plugin azure
  --vm-set-type VirtualMachineScaleSets
  --vnet-subnet-id $AKSSUBNETID
  --docker-bridge-address $AKS_NETWORKING_DOCKER_BRIDGE_ADDRESS
  --dns-service-ip $AKS_NETWORKING_DNS_SERVICE_IP
  --service-cidr $AKS_NETWORKING_SERVICE_CIDR
  --aad-server-app-id $AADSERVERAPPID
  --aad-server-app-secret $AADSERVERAPPSECRET
  --aad-client-app-id $AADCLIENTAPPID
  --aad-tenant-id $TENANTID
```





Azure resources for AKS



azureserbia

| | NAME ↑↓ | TYPE ↑↓ | LOCATION ↑↓ |
|--|---------|--------------------|-------------|
| | aksdemo | Kubernetes service | East US |

Scale

Scale method Manual [Autoscale](#)

Node count 3 x Standard B2s (2 vcpus, 4 GiB memory)

Total cluster capacity

| | |
|--------|---------|
| Cores | 6 vCPUs |
| Memory | 12 GiB |

Upgrade

Search (Ctrl+/Save Discard

You can upgrade your cluster to a newer version of Kubernetes. This will upgrade the control plane components of your cluster. To upgrade your node pools, go to the 'Node pools' menu item instead.

[Learn more about upgrading your AKS cluster](#)
[View the Kubernetes changelog](#)

Kubernetes version

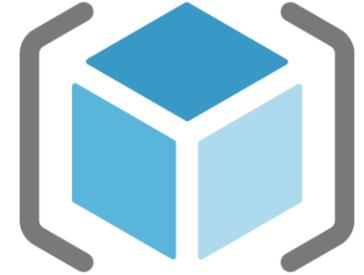
1.15.7 (c... ^)
1.16.7
1.15.10
1.15.7 (current)

Settings

Node pools
 Upgrade
 Scale



Azure resources for AKS



[cube] azureserbia

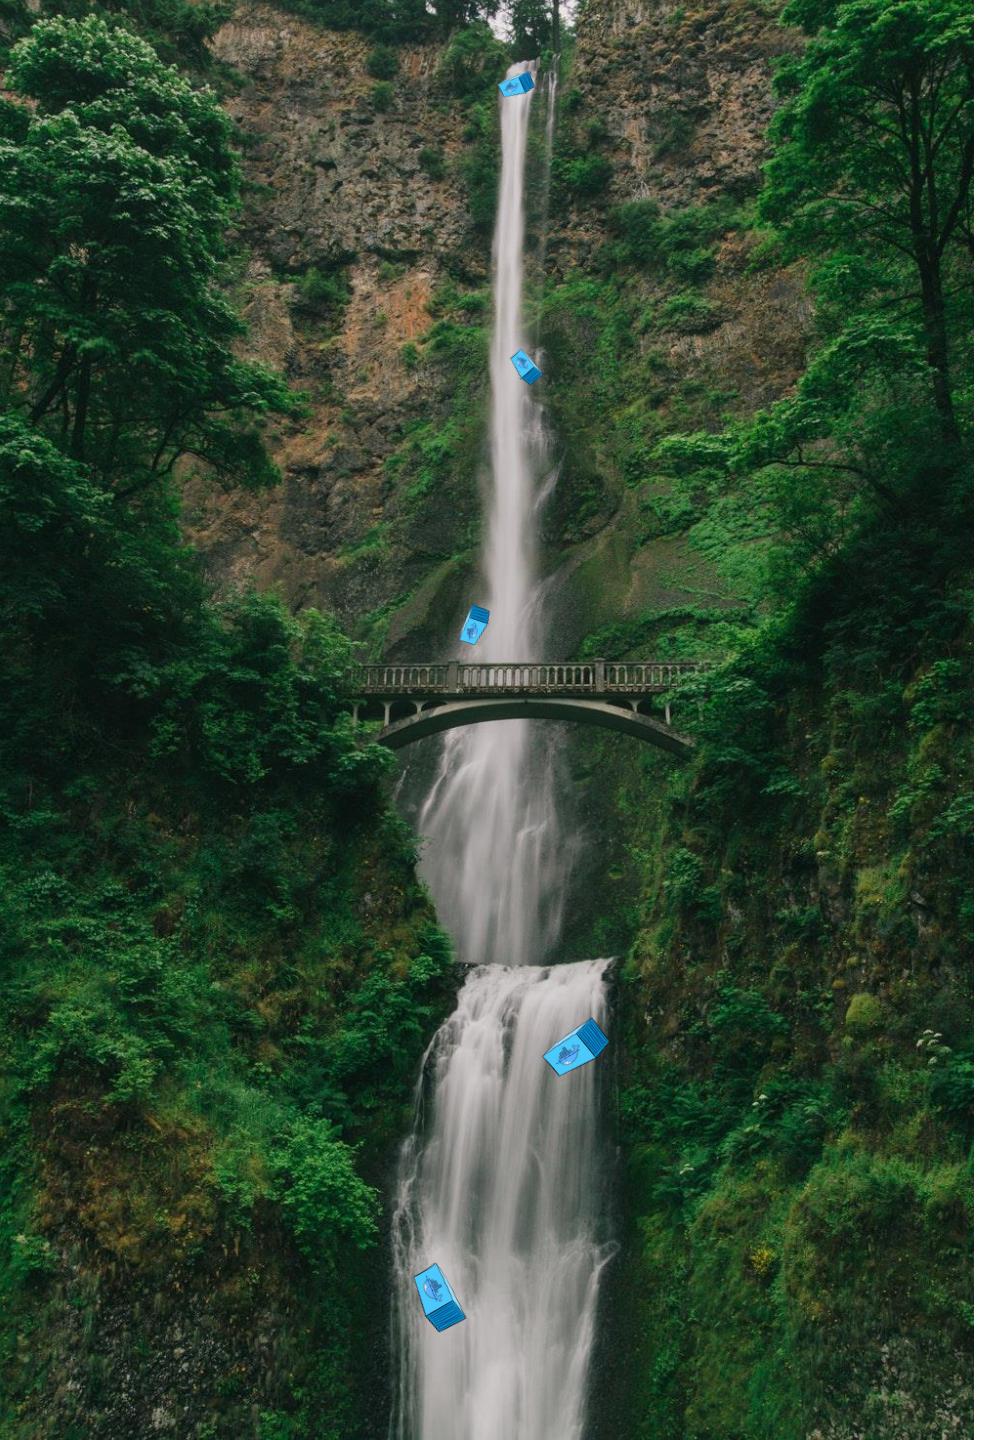
| <input type="checkbox"/> NAME ↑↓ | TYPE ↑↓ | LOCATION ↑↓ |
|----------------------------------|--------------------|-------------|
| <input type="checkbox"/> aksdemo | Kubernetes service | East US |

[cube] MC_azureserbia_aksdemo_westeurope

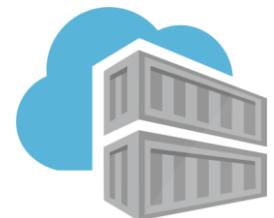
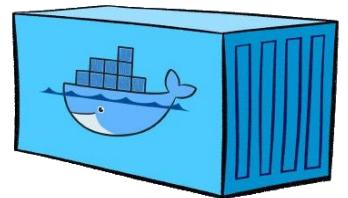
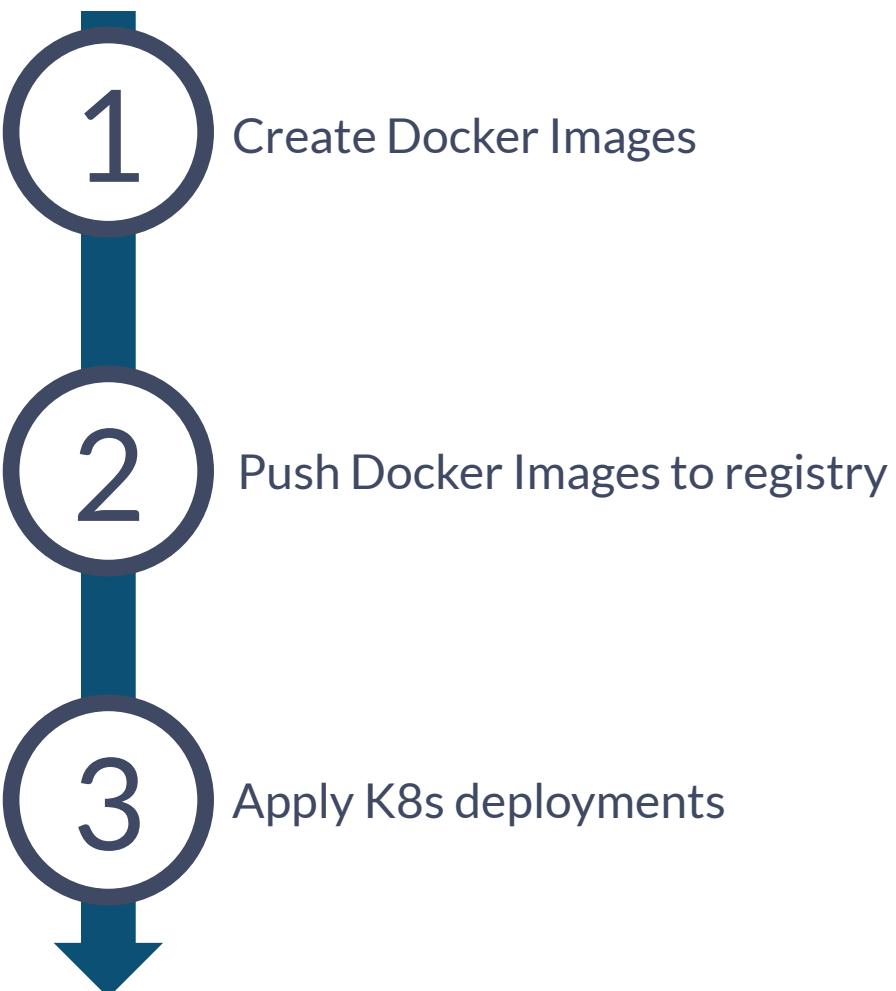
| <input type="checkbox"/> Name ↑↓ | Type ↑↓ | Location ↑↓ |
|--|---------------------------|-------------|
| <input type="checkbox"/> 4e71e966-89b3-488f-9c01-1894b33b27d2 | Public IP address | West Europe |
| <input type="checkbox"/> aks-agentpool-31287680-nsg | Network security group | West Europe |
| <input type="checkbox"/> aks-linuxpool-31287680-vmss | Virtual machine scale set | West Europe |
| <input type="checkbox"/> aks2go-aks-we-agentpool | Managed Identity | West Europe |
| <input type="checkbox"/> kubernetes | Load balancer | West Europe |
| <input type="checkbox"/> kubernetes-a1ab23421796744eaaad275b435ee5c2 | Public IP address | West Europe |
| <input type="checkbox"/> kubernetes-a93232242301947a0ae2463f322bf102 | Public IP address | West Europe |
| <input type="checkbox"/> omsagent-aks2go-aks-we | Managed Identity | West Europe |

Post Install steps

```
## download & install kubectl  
az aks install-cli  
  
## download credentials  
az aks get-credentials --resource-group communitylive --name aksdemo
```



Deployment flow



 docker hub



An aerial photograph of a winding mountain road, specifically the Stelvio Pass in the Alps, which is known for its many sharp turns and hairpin curves. The road is dark grey asphalt, contrasting with the surrounding green grass and rocky terrain. The perspective is from above, looking down the length of the road as it disappears into the distance.

Kubernetes concepts



@pascalnaber

Pod

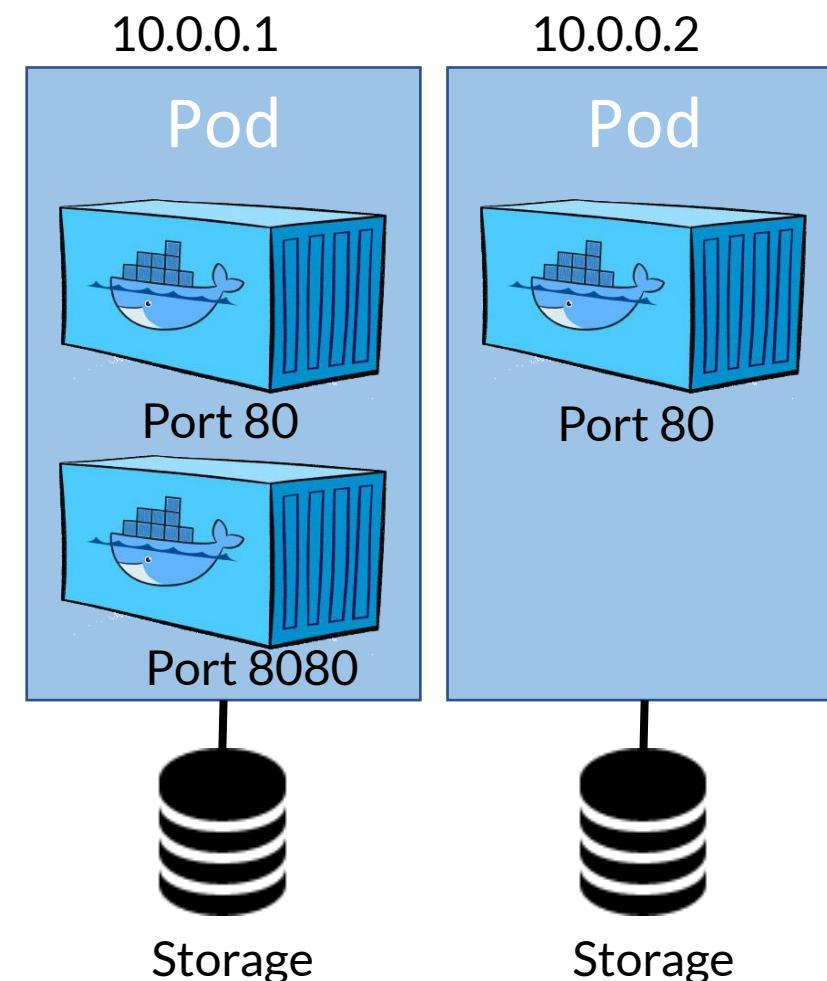
Group of 1 or more containers

Shared Storage

Shared Network

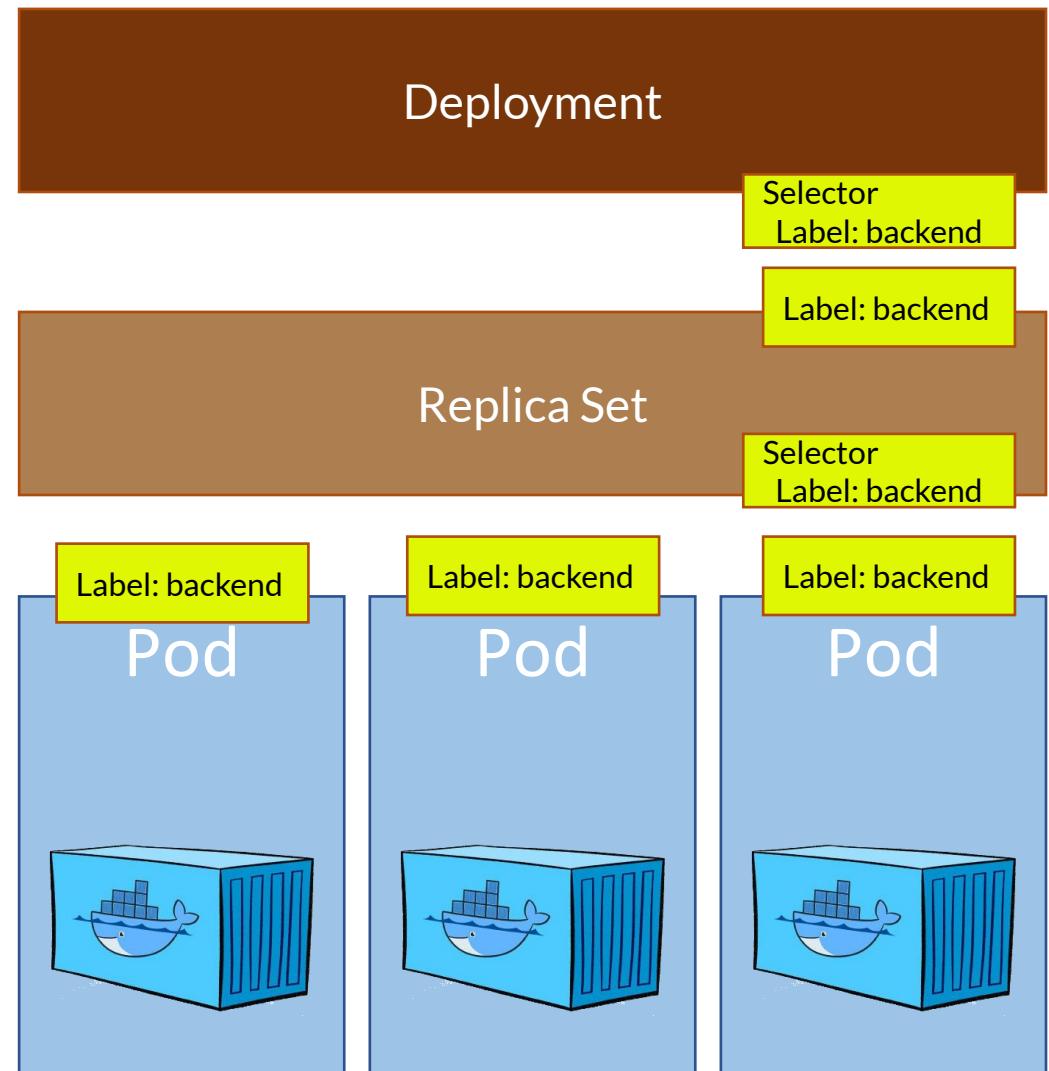
Same IP-address

Shared port-range



leaderboard-api.yaml

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: leaderboard-api
spec:
  replicas: 3
  template:
    metadata:
      labels:
        app: backend
    spec:
      containers:
        - name: leaderboardwebapi
          image: clouddemo.azurecr.io/leaderboardwebapi:455
          ports:
            - containerPort: 80
```



```
C:> kubectl apply -f leaderboard-api.yaml
```

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: leaderboard-api
spec:
  replicas: 10
  minReadySeconds: 5
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxUnavailable: 1
      maxSurge: 1
  template:
    metadata:
      labels:
        app: backend
    spec:
      containers:
        - name: leaderboardwebapi
          image: clouddemo.azurecr.io/leaderboardwebapi:457
      imagePullSecrets:
        - name: clouddemoimages
```

Rolling updates

Zero-downtime deployment

Can be rolled back

```
C:> kubectl apply -f update.yaml
```

Environment variables & Secrets

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: leaderboard-api
spec:
  replicas: 3
  template:
    metadata:
      labels:
        app: backend
    spec:
      containers:
        - name: leaderboardwebapi
          image: clouddemo.azurecr.io/leaderboardwebapi:457
          env:
            - name: Serilog_MinimumLevel_Default
              value: "Warning"
            - name: ConnectionStrings__LeaderboardContext
              valueFrom:
                secretKeyRef:
                  name: sqlserver
                  key: connectionstring
      ports:
        - containerPort: 80
  imagePullSecrets:
    - name: clouddemoimages
```

Environment variables & Secrets

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: leaderboard-api
spec:
  replicas: 3
  template:
    metadata:
      labels:
        app: backend
    spec:
      containers:
        - name: leaderboardwebapi
          image: clouddemo.azurecr.io/leaderboardwebapi:457
          env:
            - name: Serilog_MinimumLevel_Default
              value: "Warning"
            - name: ConnectionStrings__LeaderboardContext
              valueFrom:
                secretKeyRef:
                  name: sqlserver
                  key: connectionstring
        ports:
          - containerPort: 80
imagePullSecrets:
  - name: clouddemoimages
```

Secrets:

Base64 encoded values

```
apiVersion: v1
kind: Secret
metadata:
  name: sqlserver
type: Opaque
data:
  connectionstring: U2VydmVyPXRjcDpnYW1pbmU=
```

```
C:> kubectl apply -f secret.yaml
```

Create Secret to access Azure Container Registry

```
kubectl create secret docker-registry clouddemoimages
--docker-server=https://clouddemo.azurecr.io
--docker-username=clouddemo --docker-password=kD98ddl$=
--docker-email=pnaber@xpirit.com
```

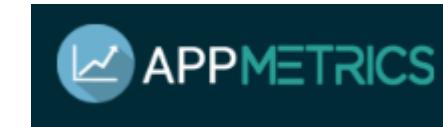
```
minReadySeconds: 5
strategy:
  type: RollingUpdate
  rollingUpdate:
    maxUnavailable: 1
    maxSurge: 1
template:
  metadata:
    labels:
      app: backend
spec:
  containers:
    - name: leaderboardwebapi
      image: clouddemo.azurecr.io/leaderboardwebapi:455
      ports:
        - containerPort: 80
livenessProbe:
  httpGet:
    path: /health
    port: 80
    initialDelaySeconds: 2
    timeoutSeconds: 1
    periodSeconds: 10
    failureThreshold: 3
readinessProbe:
  httpGet:
    path: /health
    port: 80
    initialDelaySeconds: 2
    timeoutSeconds: 1
    periodSeconds: 10
    failureThreshold: 3
```

Health checks

livenessProbe

Indicates whether the Container is running
Restart in case of failure

Container:



<https://www.app-metrics.io/>

HealthChecks library

<https://github.com/dotnet-architecture/HealthChecks>

Microsoft.AspNetCore.Diagnostics.HealthChecks

readinessProbe

Indicates whether the Container is ready to service requests
No traffic is routed to the Pod

```
C:> kubectl apply -f leaderboard-api.yaml
```

Kubernetes – Self healing

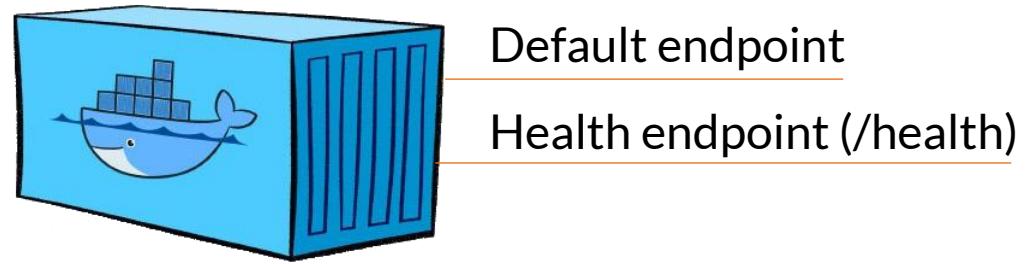
Health endpoints returns != 200?

- ⟳ Every n seconds check
Restart container

🚫 During rolling update deployment

🚫 Stop deployment

✖️ During container startup
No traffic



```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: leaderboard-api
spec:
  replicas: 10
  minReadySeconds: 5
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxUnavailable: 1
      maxSurge: 1
  template:
    metadata:
      labels:
        app: backend
    containers:
      - name: leaderboardwebapi
        image: clouddemo.azurecr.io/leaderboardwebapi:457
        resources:
          requests:
            cpu: "100m"
            memory: "64Mi"
          limits:
            cpu: "500m"
            memory: "128Mi"
  env:
    - name: Serilog__MinimumLevel__Default
      value: "Warning"
    - name: ConnectionStrings__LeaderboardContext
      valueFrom:
        secretKeyRef:
          name: sqlserver
          key: connectionstring
          path: /health
          port: 80
        initialDelaySeconds: 2
        timeoutSeconds: 1
```

Resource Management

requests: Minimum required resources
limits: Capped resource usage

100m = 0.1 cpu

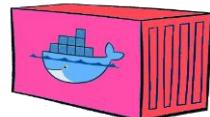
500m = 0.5 cpu

64Mi = 64 MB memory

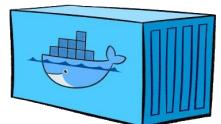
128Mi = 128 MB memory

```
C:> kubectl apply -f leaderboard-api.yaml
```

Azure Kubernetes Service (AKS) - Scaling



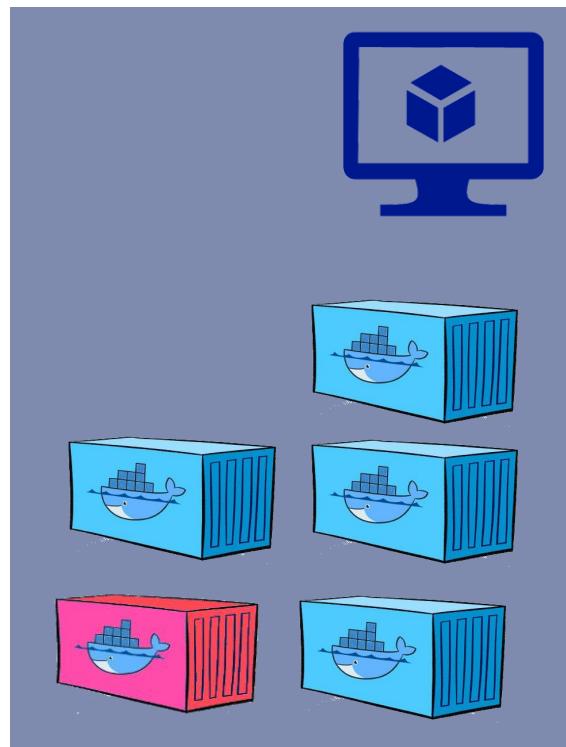
5 replicas



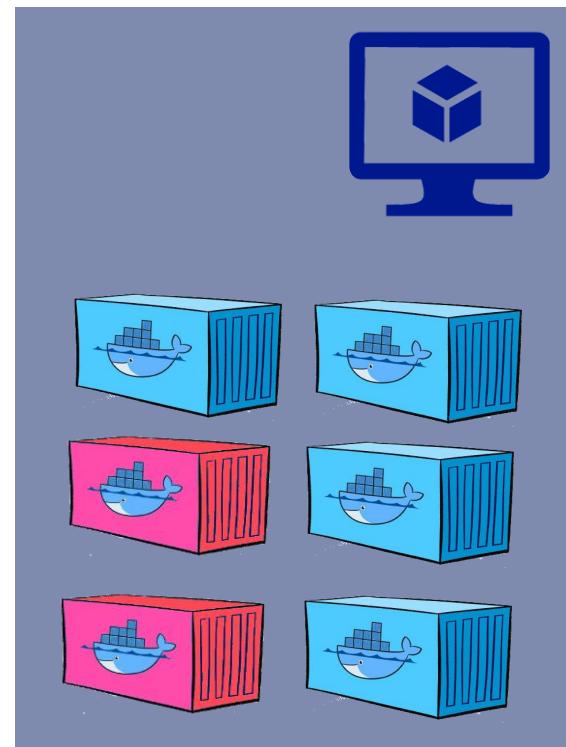
4 replicas

Pod Autoscaler 4-20
> 60% CPU

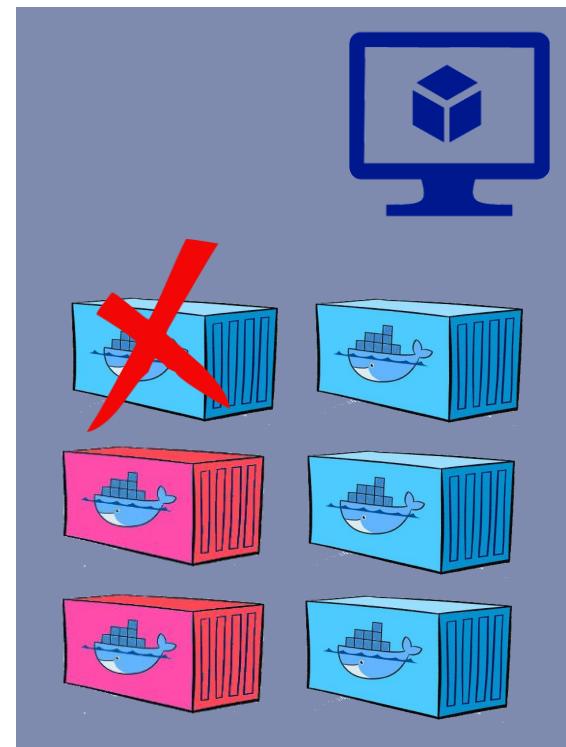
Cluster Autoscaler



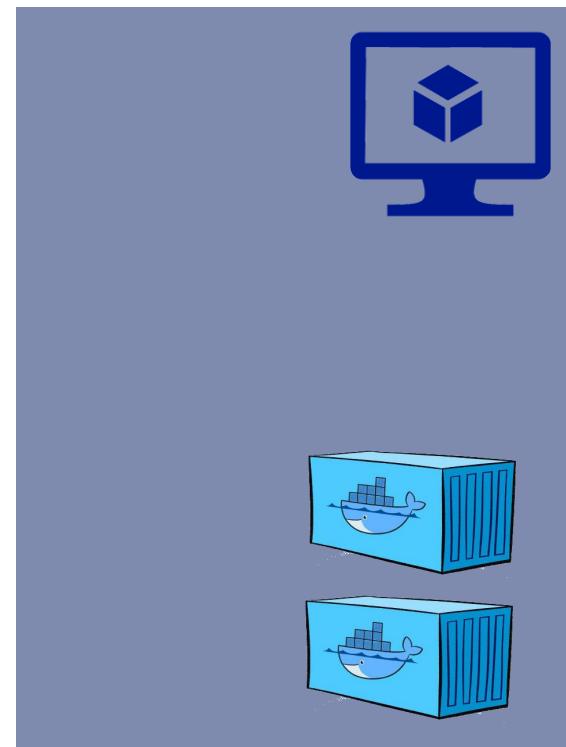
worker



worker



worker



worker



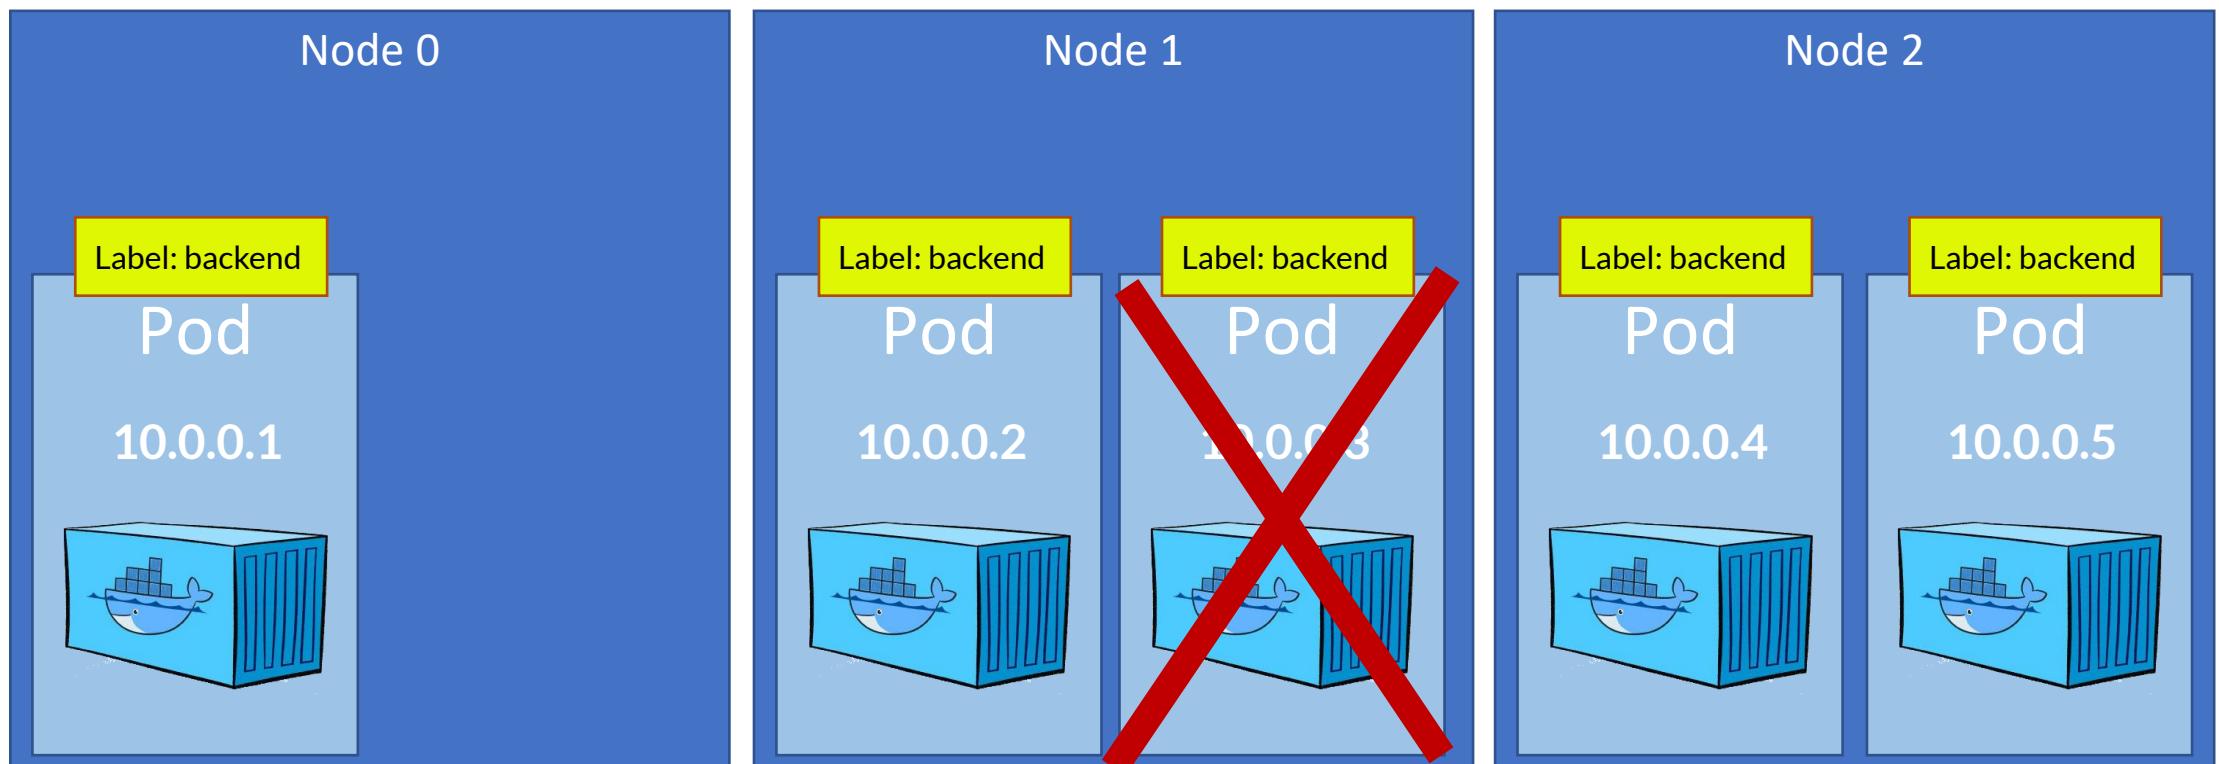
How to access the Pods?

From inside and outside the cluster

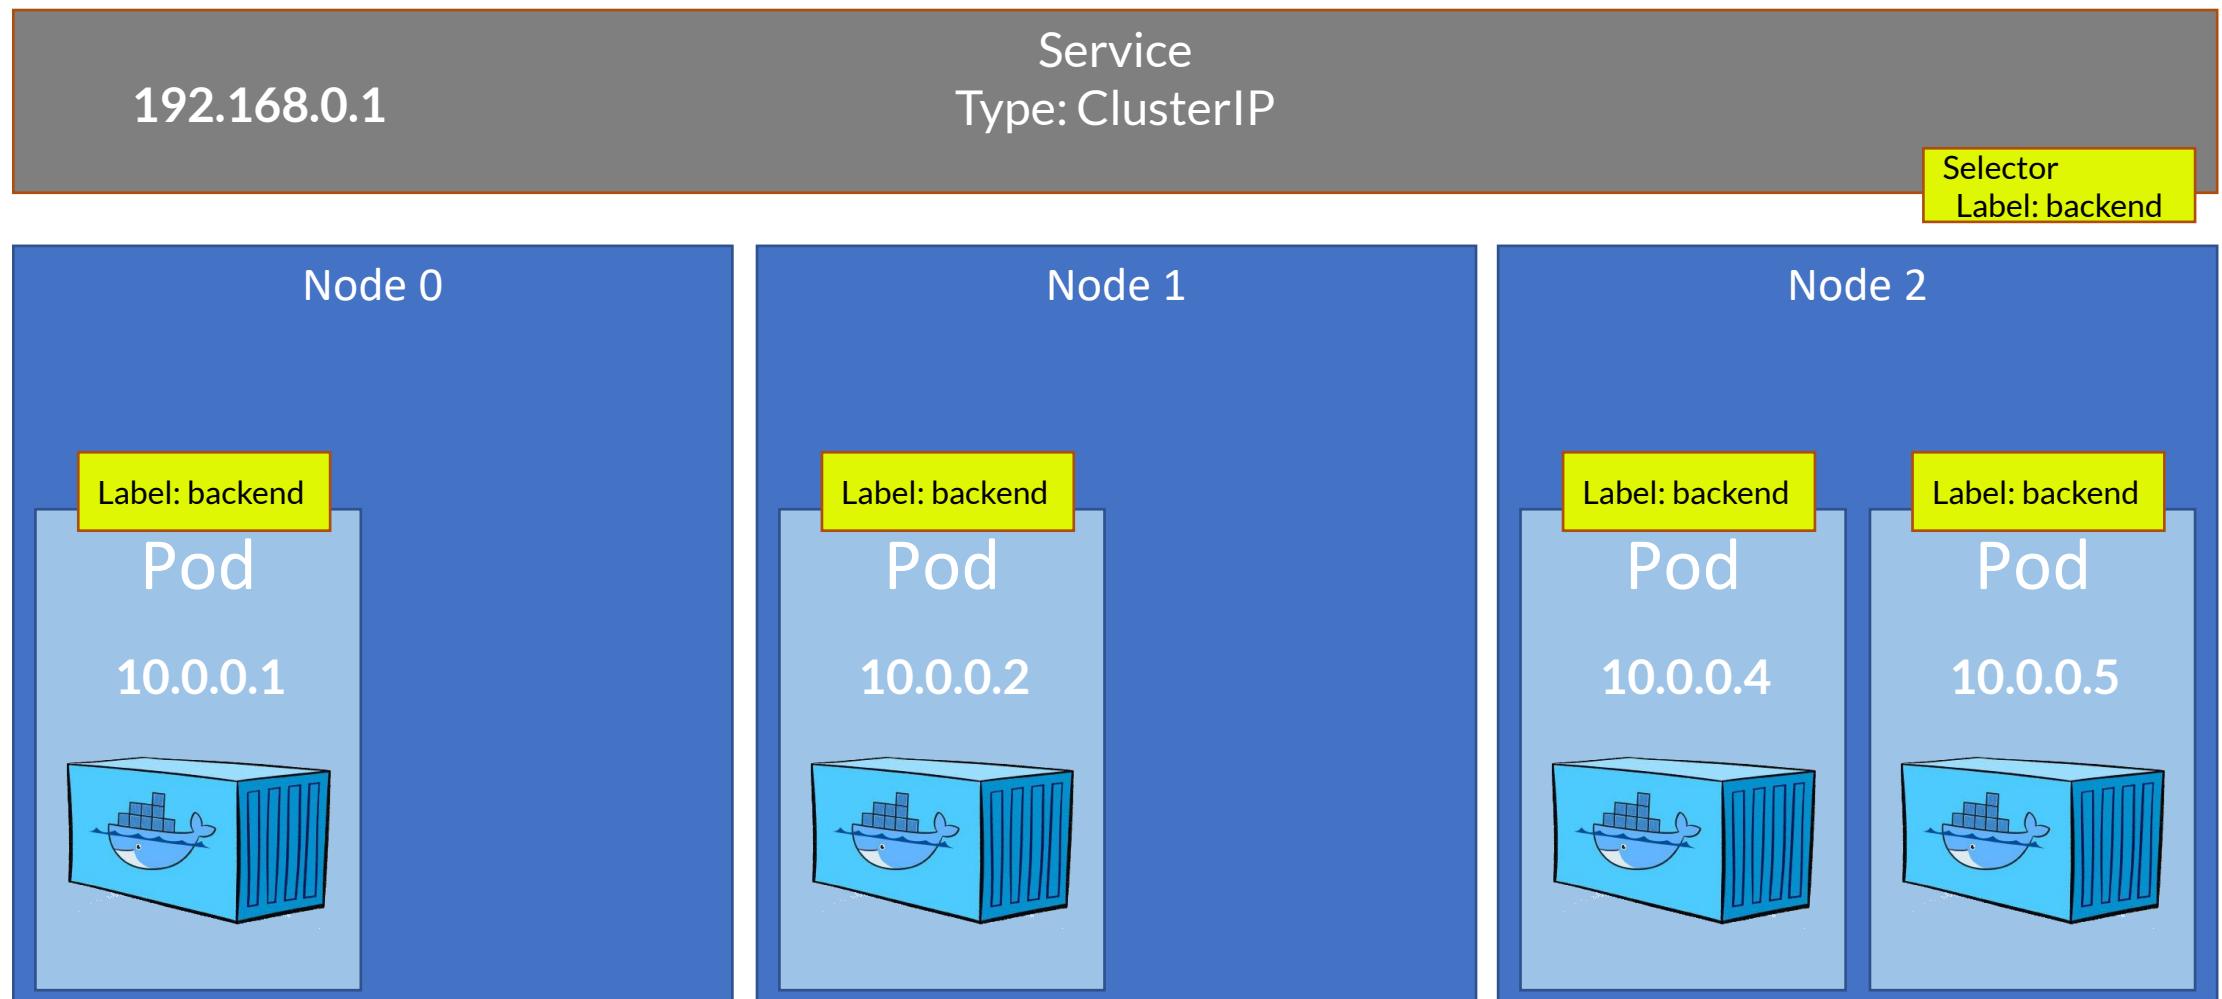


@pascalnaber

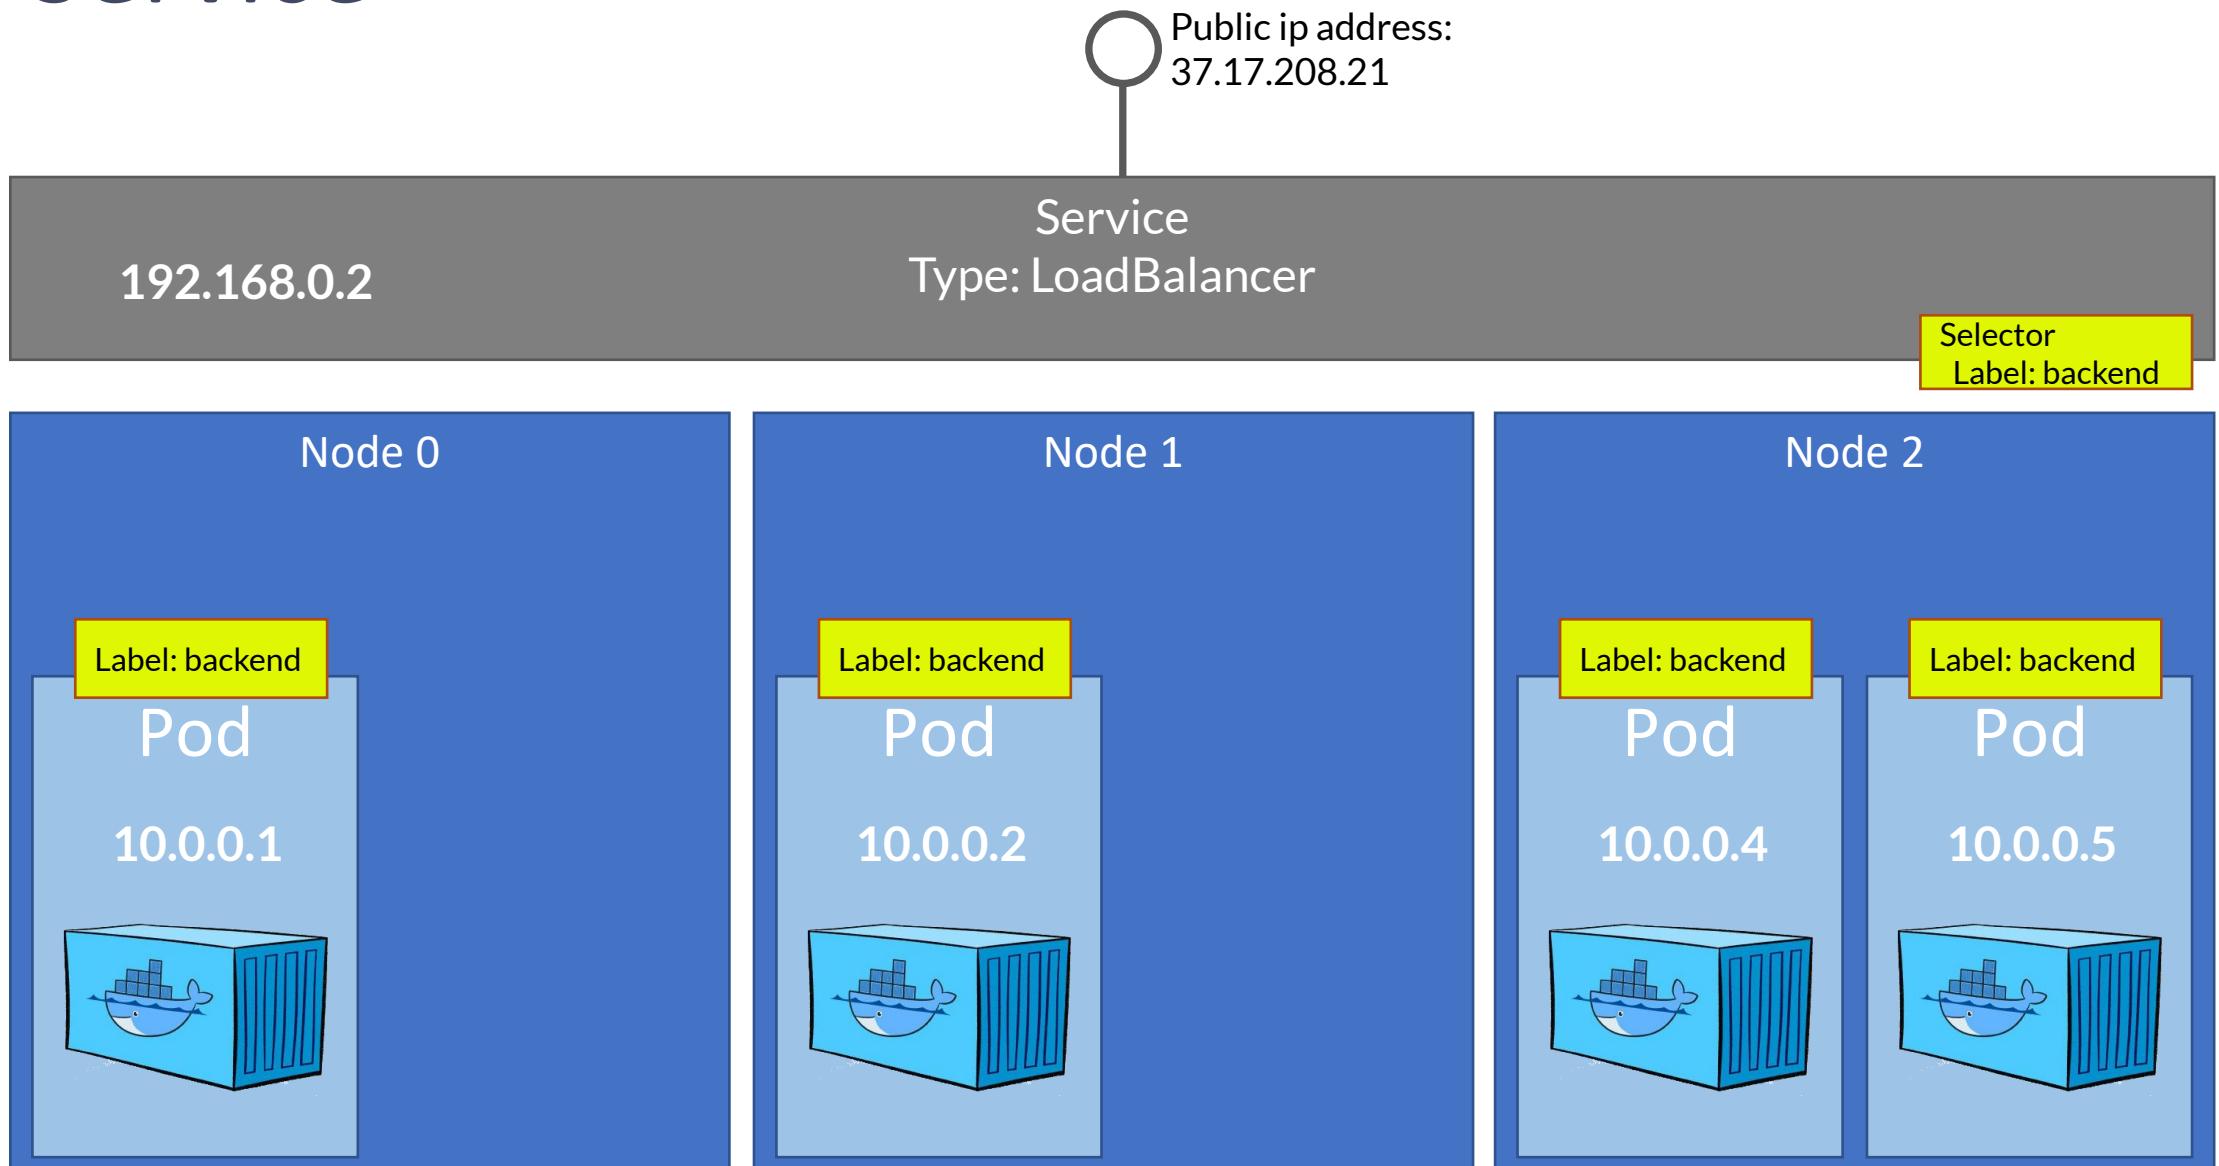
Pods are mortal

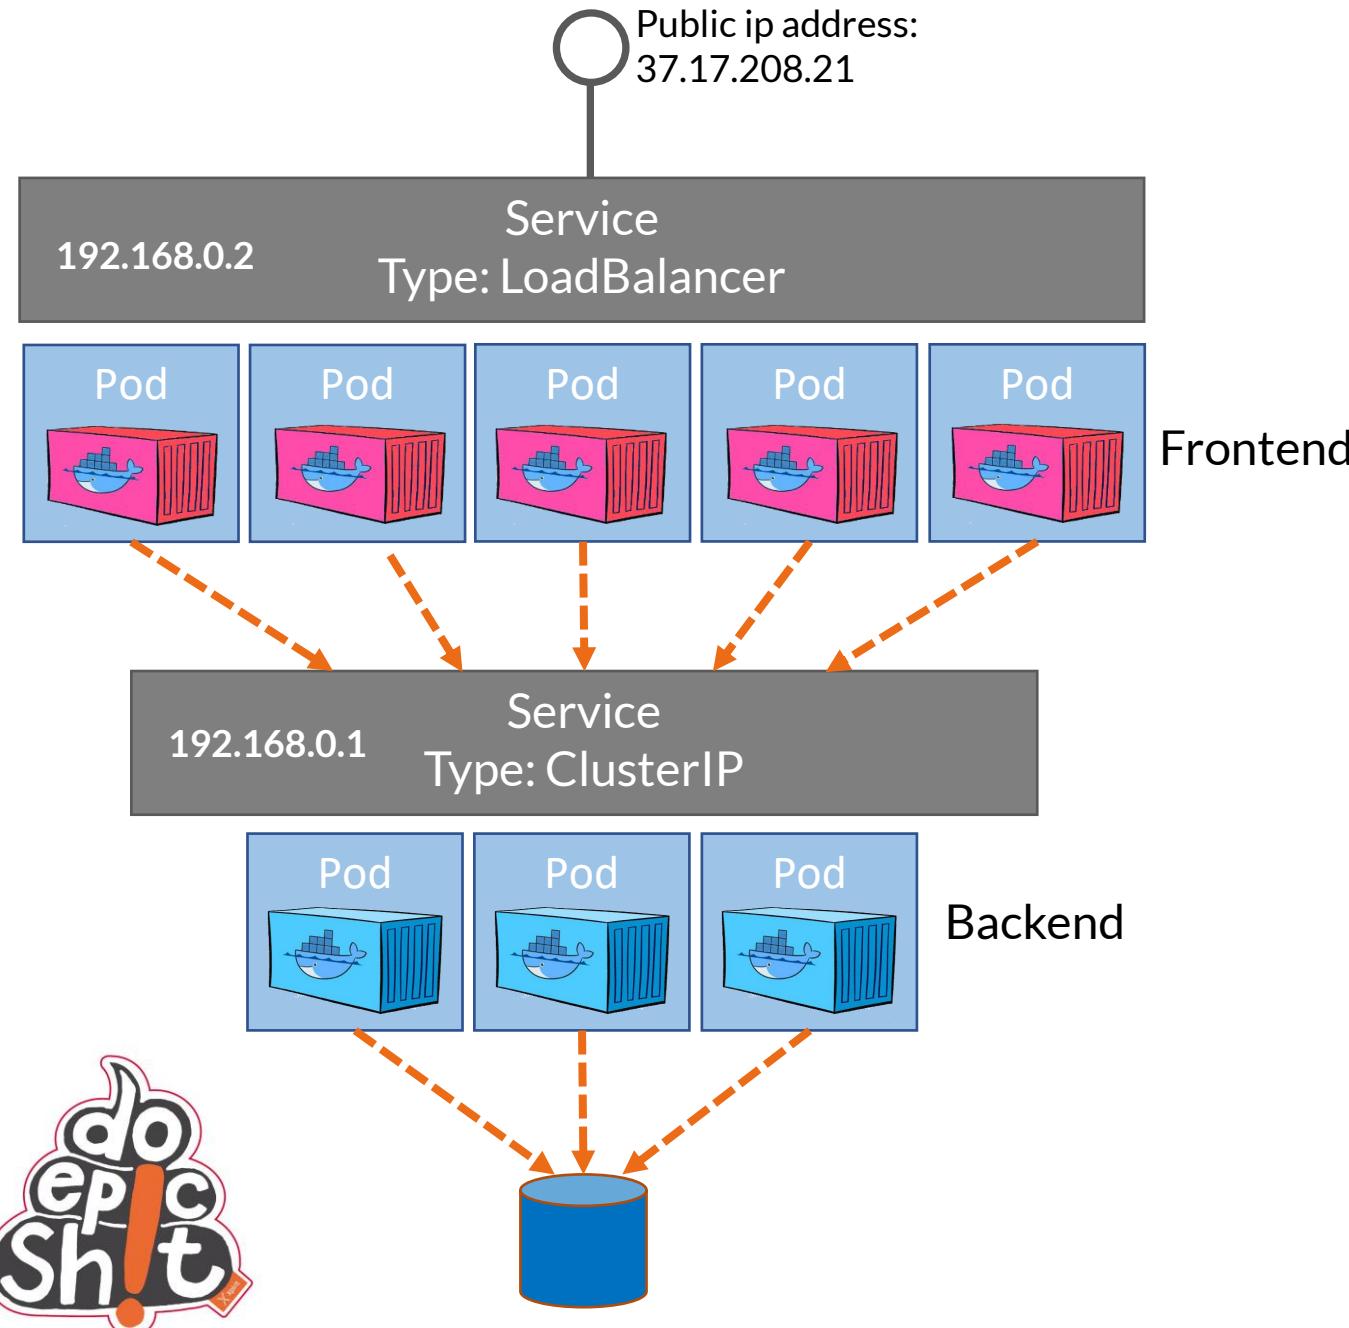


Service



Service





```

apiVersion: v1
kind: Service
metadata:
  name: leaderboardapi-service
  namespace: service
spec:
  ports:
    - port: 80
  selector:
    app: backend
  type: ClusterIP

```

```
C:> kubectl apply -f leaderboard-service.yaml
```

```

spec:
  containers:
    - name: gamingwebapp
      image: clouddemo.azurecr.io/gamingwebapp:455
      env:
        - name: LeaderboardWebApiBaseUrl
          value: "http://leaderboardapi-service"
      ports:
        - containerPort: 80

```

```
C:> kubectl apply -f gamingwebapp.yaml
```

Too many ip-addresses. Now what?

I want to access all my services through the same ip-address



@pascalnaber

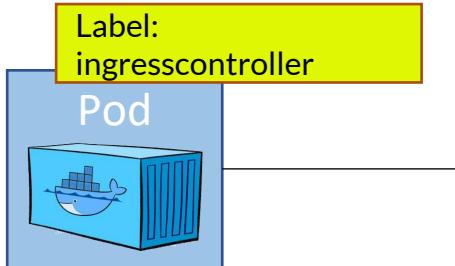
Ingress



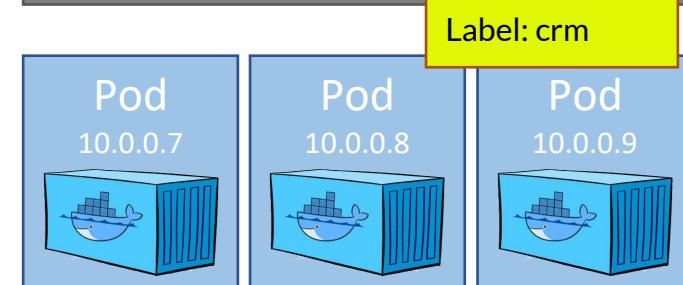
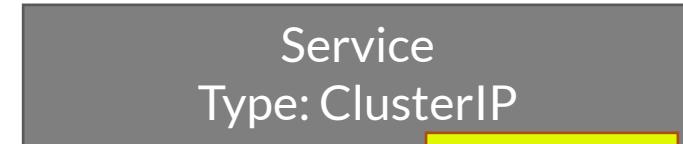
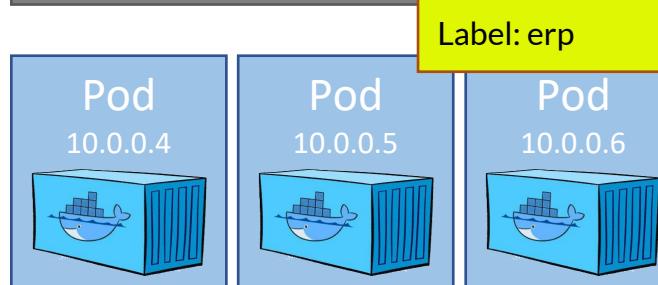
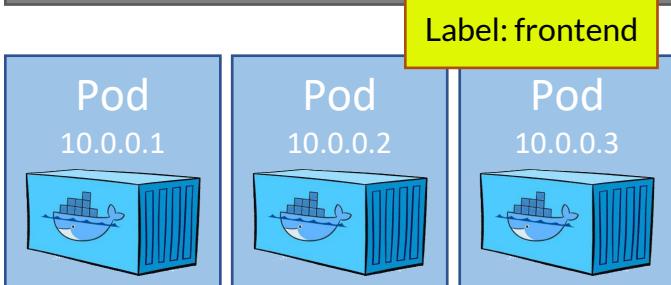
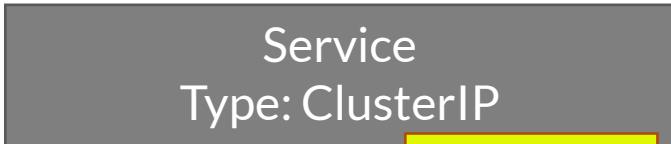
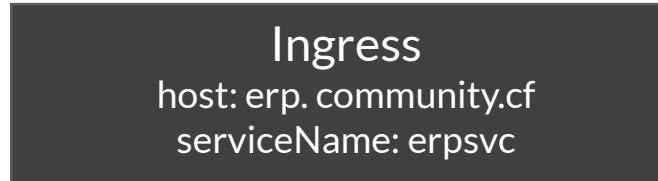
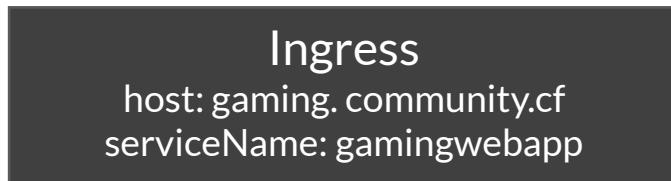
Public ip address:
37.17.208.21



quay.io/kubernetes-ingress-controller/
nginx-ingress-controller:0.15.0



```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: gamingwebapp-ingresscf
  namespace: ingressdemo
  annotations:
    kubernetes.io/ingress.class: nginx
spec:
  rules:
  - host: gaming.matrixproject.cf
    http:
      paths:
      - backend:
          serviceName: gamingwebapp-service
          servicePort: 80
```





Deployment
Replica set
Pod
Label
Rolling update
Health check
Environment variables
Secret
Resource management
Horizontal Pod Autoscaler

Namespace
Service
Ingress
Annotation
Persistent Volume
Cron Job
Daemon Set
Job
Stateful Set
Config Map



@pascalnaber

Azure Dev Spaces



Cert manager



Ecosystem



= Kubeflow





Helm

Helm is a tool for managing
packages of pre-configured
Kubernetes resources

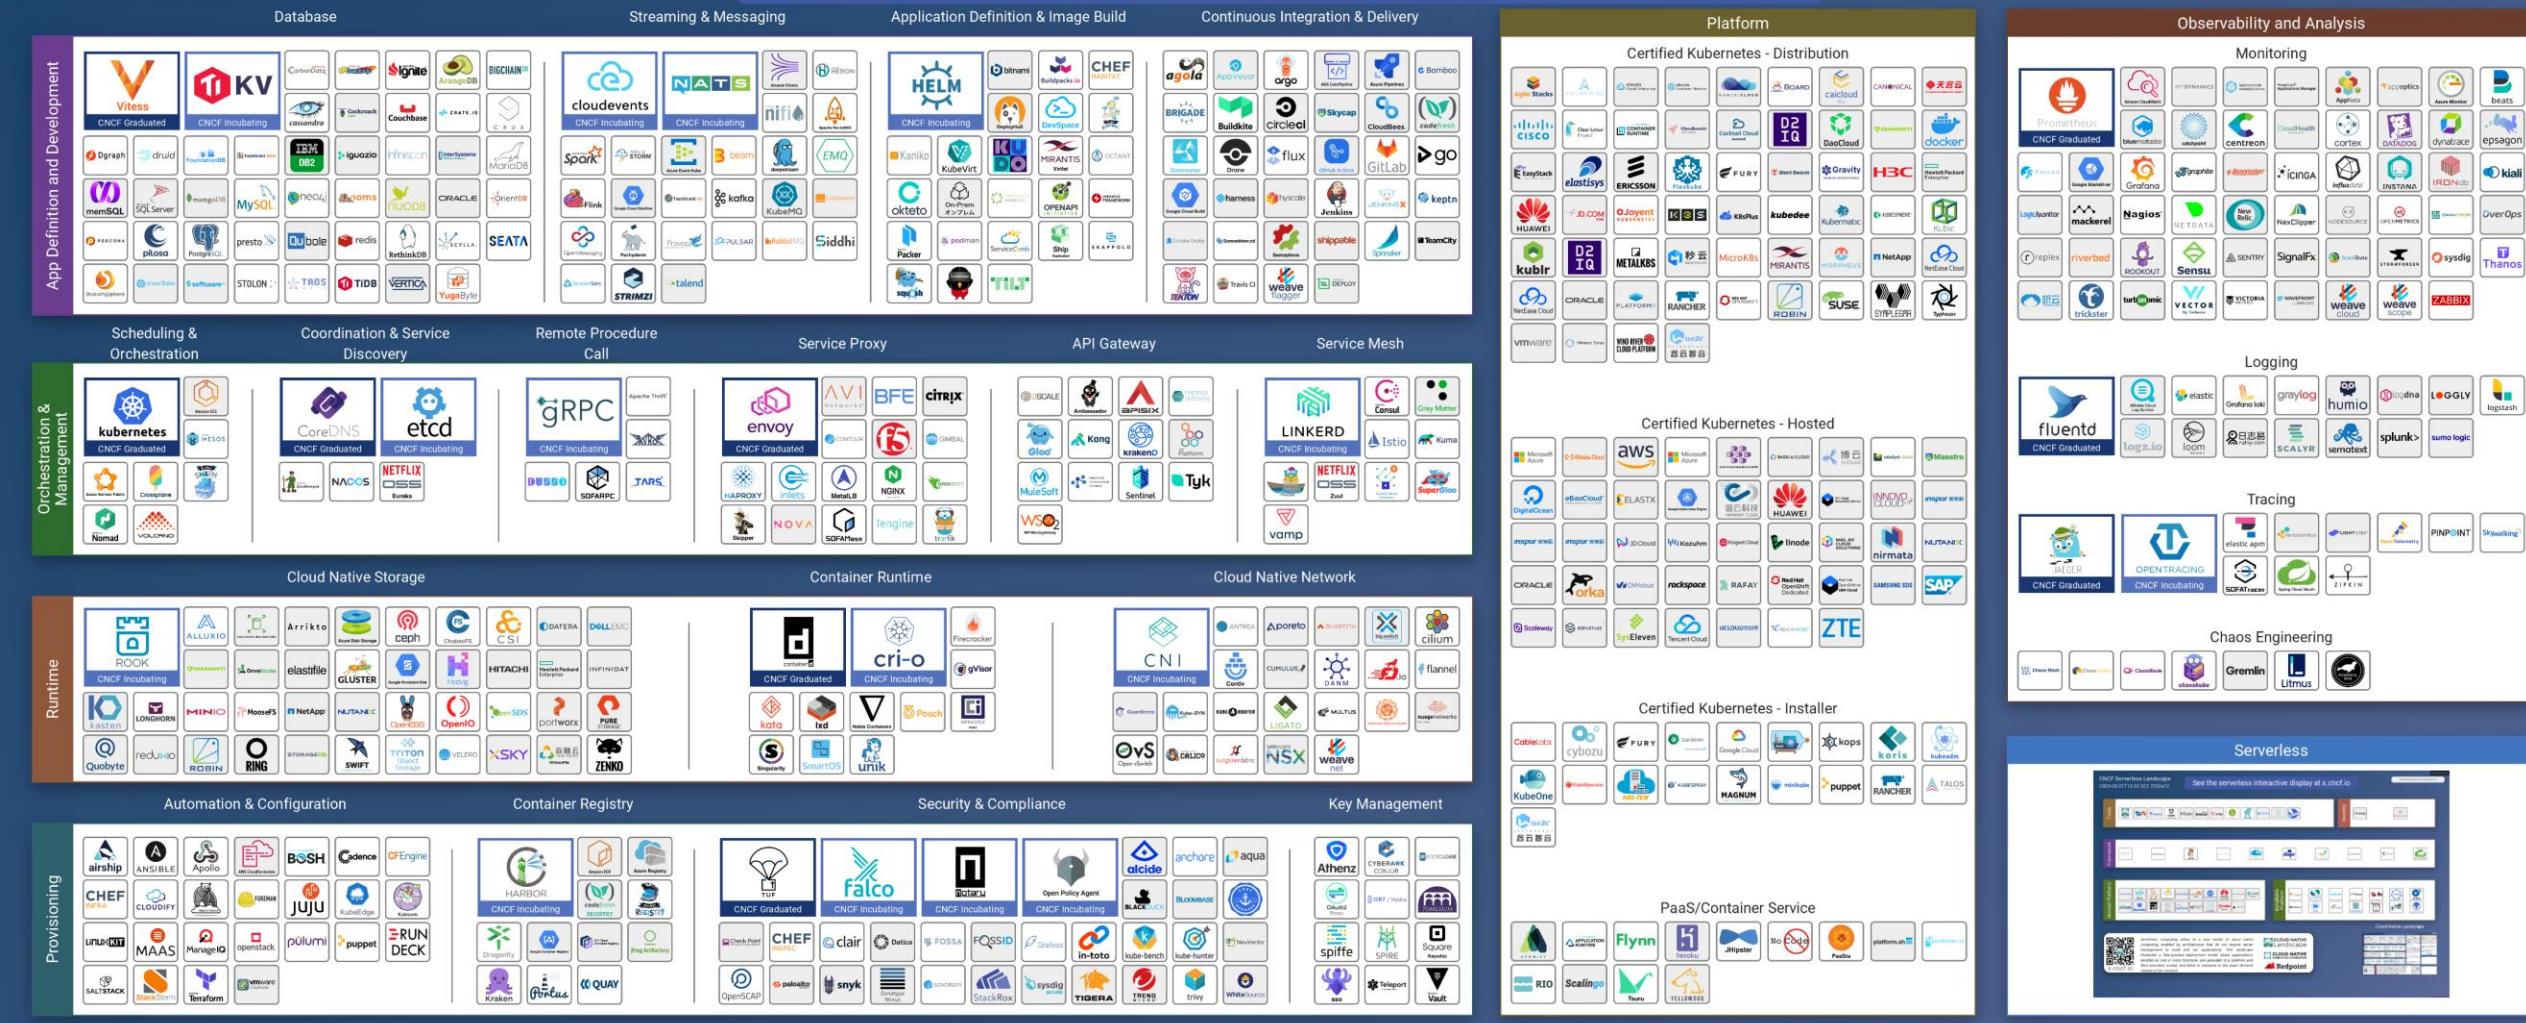
<https://github.com/kubernetes/helm>

Artifactory
Cert-manager
Consul
Cassandra
Couchdb
Datadog
Docker-registry
Drupal
Elasticsearch
Ethereum
fluentd
Gitlab
Grafana
Hadoop
Heapster
Jenkins
Joomla
Kafka
Kibana
Kong
Kubeless
Lamp

Linkerd
Logstash
Magento
Mariadb
Mongodb
Mssql-linux
Mysql
Neo4j
Newrelic
Openvpn
Postgresql
Presto
Prometheus
Rabbitmq
Redis
Selenium
Sonarqube
Spinnaker
Sysdig
Tensorflow
Traefik
Wordpress

Overwhelmed? Please see the CNCF Trail Map. That and the interactive landscape are at l.cncf.io

Greyed logos are not open source



Kubernetes Ecosystem

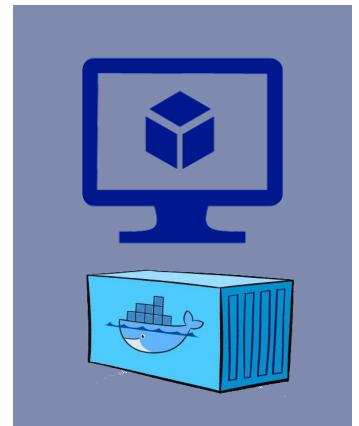
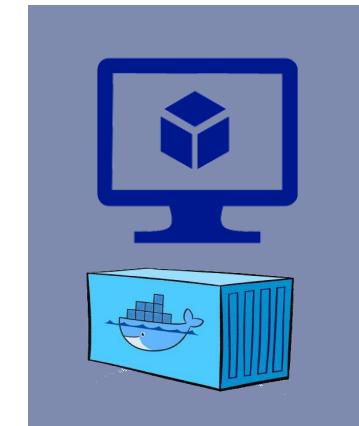
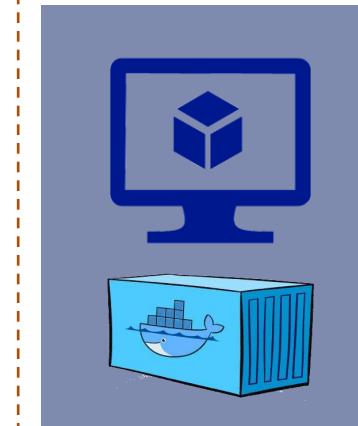
Azure Kubernetes Service (AKS)



100% managed by Microsoft



IaaS managed by Microsoft



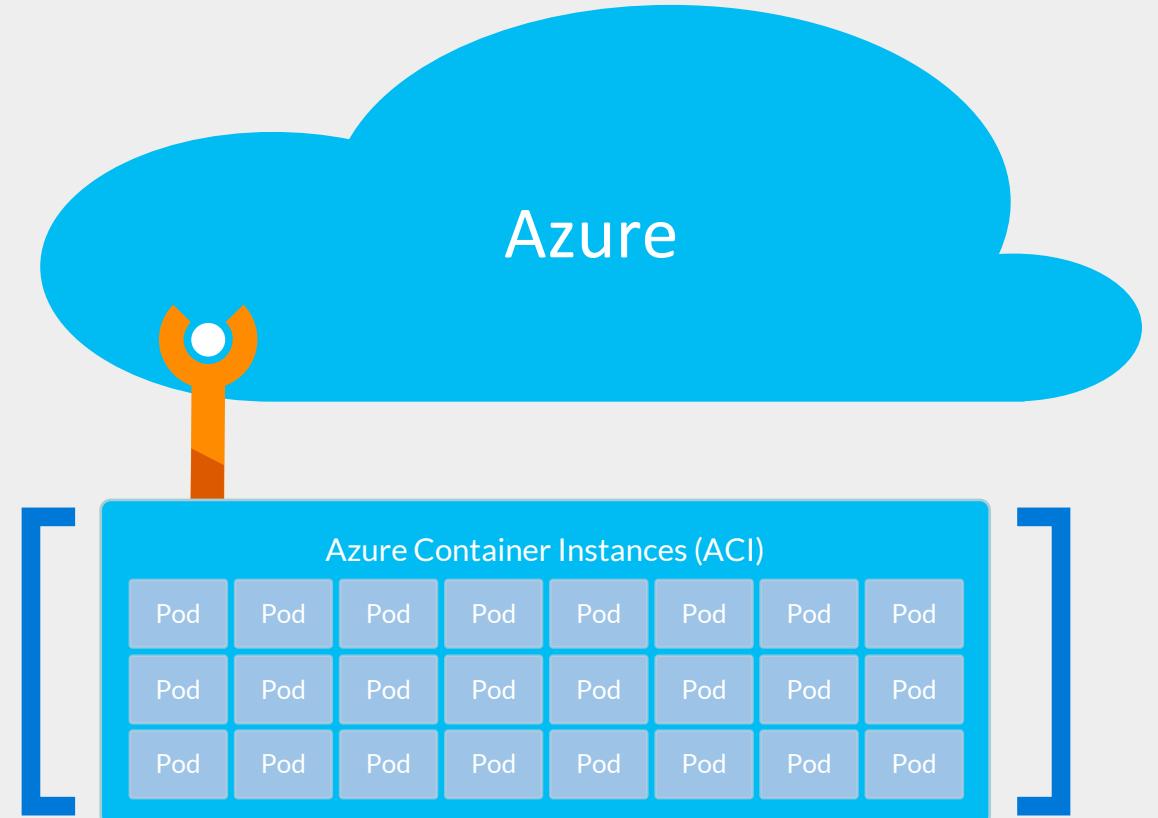
Azure Container Instances (ACI)

Starts in seconds

No VM Management

Billed per second

Linux and Windows containers



AKS + ACI



Virtual
Kubelet

100% managed by Microsoft



master

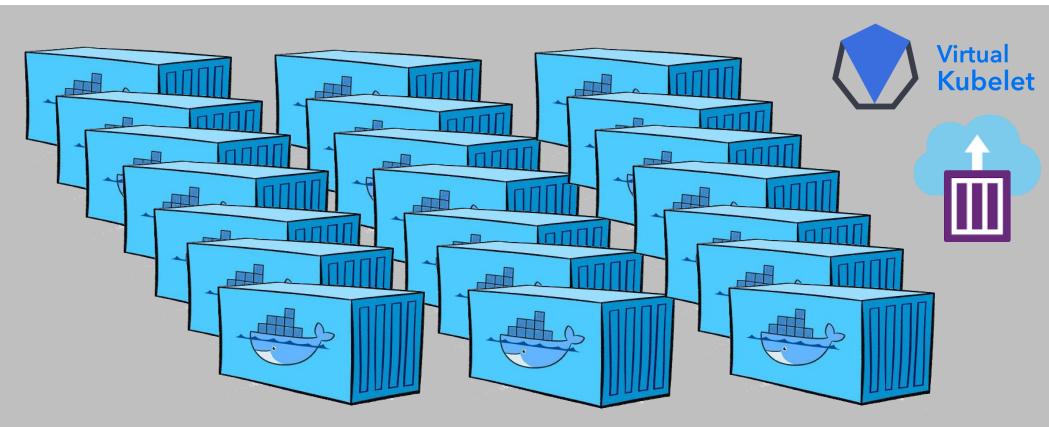


master



master

100% managed by Microsoft



worker



AKS





The future of
infrastructure is
happening now



kubernetes

 @pascalnaber

Wrap up

Docker

Install applications to a container instead of a server

Ship containers everywhere

Kubernetes in the cloud

Deploy your containerized workloads on Kubernetes

Brings a clean separation between Dev and Ops

Provides scaling, rolling updates, high availability





P4l Naber

Coding Azure Architect

Xpirit Netherlands

@pascalnaber

<http://pascalnaber.wordpress.com>

<https://github.com/pascalnaber/azureserbia>

